

CEE 616: Probabilistic Machine Learning Foundations: Optimization

Jimi Oke

UMassAmherst

College of Engineering

Thu, Sep 18, 2025

Outline

- ① Introduction
- ② First-order methods
- ③ Second-order methods
- ④ Application: MLE
- ⑤ Constrained optimization
- ⑥ Outlook


Optimization

Optimization

Optimization is the body of mathematics that deals with the theory and algorithms for characterizing the maximum/minimum values of functions.


Optimization

Optimization is the body of mathematics that deals with the theory and algorithms for characterizing the maximum/minimum values of functions.



Optimization


Optimization is the body of mathematics that deals with the theory and algorithms for characterizing the maximum/minimum values of functions.



We will consider two widely-used approaches in machine learning:

Optimization

Optimization is the body of mathematics that deals with the theory and algorithms for characterizing the maximum/minimum values of functions.




We will consider two widely-used approaches in machine learning:

- First-order methods (e.g. gradient descent)

Optimization

Optimization is the body of mathematics that deals with the theory and algorithms for characterizing the maximum/minimum values of functions.



We will consider two widely-used approaches in machine learning:

- First-order methods (e.g. gradient descent)
- Second-order methods (e.g. Newton's method)

Definitions

Definitions

- Minimum of convex function

Definitions

- Minimum of convex function

Definitions

- Minimum of convex function

$$\theta^* =$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) :$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) =$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

- Gradient of a function

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

- Gradient of a function

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

- Gradient of a function (vector of partial derivatives):

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

- Gradient of a function (vector of partial derivatives):

$$\nabla_{\theta} \mathcal{L}(\theta) =$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

- Gradient of a function (vector of partial derivatives):

$$\nabla_{\theta} \mathcal{L}(\theta) = \begin{bmatrix} \frac{\partial \mathcal{L}(\theta)}{\partial \theta_1} \end{bmatrix}$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

- Gradient of a function (vector of partial derivatives):

$$\nabla_{\theta} \mathcal{L}(\theta) = \begin{bmatrix} \frac{\partial \mathcal{L}(\theta)}{\partial \theta_1} \\ \frac{\partial \mathcal{L}(\theta)}{\partial \theta_2} \end{bmatrix}$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

- Gradient of a function (vector of partial derivatives):

$$\nabla_{\theta} \mathcal{L}(\theta) = \begin{bmatrix} \frac{\partial \mathcal{L}(\theta)}{\partial \theta_1} \\ \frac{\partial \mathcal{L}(\theta)}{\partial \theta_2} \\ \vdots \end{bmatrix}$$

Definitions

- Minimum of convex function

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta) : \frac{d\mathcal{L}(\theta^*)}{d\theta} = 0 \quad (1)$$

- Derivative of a function:

$$\mathcal{L}'(\theta) = \frac{d\mathcal{L}(\theta)}{d\theta} \quad (2)$$

- Gradient of a function (vector of partial derivatives):

$$\nabla_{\theta} \mathcal{L}(\theta) = \begin{bmatrix} \frac{\partial \mathcal{L}(\theta)}{\partial \theta_1} \\ \frac{\partial \mathcal{L}(\theta)}{\partial \theta_2} \\ \vdots \\ \frac{\partial \mathcal{L}(\theta)}{\partial \theta_n} \end{bmatrix} \quad (3)$$

Further definitions

Further definitions

- Second-order derivative:

Further definitions

- Second-order derivative:

$$\mathcal{L}''(\theta) =$$

Further definitions

- Second-order derivative:

$$\mathcal{L}''(\theta) = \frac{d^2 \mathcal{L}(\theta)}{d\theta^2} \quad (4)$$

Further definitions

- Second-order derivative:

$$\mathcal{L}''(\theta) = \frac{d^2 \mathcal{L}(\theta)}{d\theta^2} \quad (4)$$

- Hessian (matrix of partial second derivatives):

Further definitions

- Second-order derivative:

$$\mathcal{L}''(\theta) = \frac{d^2 \mathcal{L}(\theta)}{d\theta^2} \quad (4)$$

- Hessian (matrix of partial second derivatives):

$$\mathbf{H}_{\theta}(\mathcal{L}(\theta)) = \nabla_{\theta}^2 \mathcal{L}(\theta)$$

Further definitions

- Second-order derivative:

$$\mathcal{L}''(\theta) = \frac{d^2 \mathcal{L}(\theta)}{d\theta^2} \quad (4)$$

- Hessian (matrix of partial second derivatives):

$$\mathbf{H}_{\theta}(\mathcal{L}(\theta)) = \nabla_{\theta}^2 \mathcal{L}(\theta) = \begin{bmatrix} \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_1^2} & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_1 \partial \theta_2} & \cdots & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_1 \partial \theta_n} \\ \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_2 \partial \theta_1} & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_2^2} & \cdots & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_2 \partial \theta_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_n \partial \theta_1} & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_n \partial \theta_2} & \cdots & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_n^2} \end{bmatrix} \quad (5)$$

Further definitions

- Second-order derivative:

$$\mathcal{L}''(\theta) = \frac{d^2 \mathcal{L}(\theta)}{d\theta^2} \quad (4)$$

- Hessian (matrix of partial second derivatives):

$$\mathbf{H}_{\theta}(\mathcal{L}(\theta)) = \nabla_{\theta}^2 \mathcal{L}(\theta) = \begin{bmatrix} \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_1^2} & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_1 \partial \theta_2} & \cdots & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_1 \partial \theta_n} \\ \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_2 \partial \theta_1} & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_2^2} & \cdots & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_2 \partial \theta_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_n \partial \theta_1} & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_n \partial \theta_2} & \cdots & \frac{\partial^2 \mathcal{L}(\theta)}{\partial \theta_n^2} \end{bmatrix} \quad (5)$$

Optimality conditions

Optimality conditions

For continuous, twice-differentiable functions:

Optimality conditions

For continuous, twice-differentiable functions:

- If θ^* is a local minimum, then $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* must be psd (necessary condition)

Optimality conditions

For continuous, twice-differentiable functions:

- If θ^* is a local minimum, then $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* must be psd (necessary condition)
- If $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* is pd, then θ^* is a local optimum (sufficient condition)

Optimality conditions

For continuous, twice-differentiable functions:

- If θ^* is a local minimum, then $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* must be psd (necessary condition)
- If $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* is pd, then θ^* is a local optimum (sufficient condition)

General approach

Optimality conditions

For continuous, twice-differentiable functions:

- If θ^* is a local minimum, then $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* must be psd (necessary condition)
- If $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* is pd, then θ^* is a local optimum (sufficient condition)

General approach

- Begin with an initial value θ_0

Optimality conditions

For continuous, twice-differentiable functions:

- If θ^* is a local minimum, then $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* must be psd (necessary condition)
- If $\mathbf{g}^* = \mathbf{0}$ and \mathbf{H}^* is pd, then θ^* is a local optimum (sufficient condition)

General approach

- Begin with an initial value θ_0
- At each iteration t , update θ_{t+1}
- Terminate when $\mathcal{L}(\theta_{t+1}) - \mathcal{L}(\theta_t) = \epsilon$

Convex optimization

Convex optimization

If the objective is convex and defined over a convex set, then every local minimum is a global minimum. pe

- Convex set: For any $\mathbf{x}, \mathbf{x}' \in \mathcal{S}$:

Convex optimization

If the objective is convex and defined over a convex set, then every local minimum is a global minimum. pe

- Convex set: For any $\mathbf{x}, \mathbf{x}' \in \mathcal{S}$:

$$\lambda \mathbf{x} + (1 - \lambda) \mathbf{x}' \in \mathcal{S}, \quad \forall \lambda \in [0, 1] \quad (6)$$

Convex optimization

If the objective is convex and defined over a convex set, then every local minimum is a global minimum. pe

- Convex set: For any $\mathbf{x}, \mathbf{x}' \in \mathcal{S}$:

$$\lambda \mathbf{x} + (1 - \lambda) \mathbf{x}' \in \mathcal{S}, \quad \forall \lambda \in [0, 1] \quad (6)$$

- Convex function: For any $\mathbf{x}, \mathbf{y} \in \mathcal{S}$ and for any $0 \leq \lambda \leq 1$:

Convex optimization

If the objective is convex and defined over a convex set, then every local minimum is a global minimum. pe

- Convex set: For any $\mathbf{x}, \mathbf{x}' \in \mathcal{S}$:

$$\lambda \mathbf{x} + (1 - \lambda) \mathbf{x}' \in \mathcal{S}, \quad \forall \lambda \in [0, 1] \quad (6)$$

- Convex function: For any $\mathbf{x}, \mathbf{y} \in \mathcal{S}$ and for any $0 \leq \lambda \leq 1$:

$$f(\lambda \mathbf{x} + (1 - \lambda) \mathbf{y}) \leq \lambda f(\mathbf{x}) + (1 - \lambda) f(\mathbf{y}) \quad (7)$$

Convex optimization

If the objective is convex and defined over a convex set, then every local minimum is a global minimum. pe

- Convex set: For any $\mathbf{x}, \mathbf{x}' \in \mathcal{S}$:

$$\lambda \mathbf{x} + (1 - \lambda) \mathbf{x}' \in \mathcal{S}, \quad \forall \lambda \in [0, 1] \quad (6)$$

- Convex function: For any $\mathbf{x}, \mathbf{y} \in \mathcal{S}$ and for any $0 \leq \lambda \leq 1$:

$$f(\lambda \mathbf{x} + (1 - \lambda) \mathbf{y}) \leq \lambda f(\mathbf{x}) + (1 - \lambda) f(\mathbf{y}) \quad (7)$$

Quadratic form

Given $f(\mathbf{x}) = \mathbf{x}^\top \mathbf{A} \mathbf{x}$:

Convex optimization

If the objective is convex and defined over a convex set, then every local minimum is a global minimum. pe

- Convex set: For any $\mathbf{x}, \mathbf{x}' \in \mathcal{S}$:

$$\lambda \mathbf{x} + (1 - \lambda) \mathbf{x}' \in \mathcal{S}, \quad \forall \lambda \in [0, 1] \quad (6)$$

- Convex function: For any $\mathbf{x}, \mathbf{y} \in \mathcal{S}$ and for any $0 \leq \lambda \leq 1$:

$$f(\lambda \mathbf{x} + (1 - \lambda) \mathbf{y}) \leq \lambda f(\mathbf{x}) + (1 - \lambda) f(\mathbf{y}) \quad (7)$$

Quadratic form

Given $f(\mathbf{x}) = \mathbf{x}^\top \mathbf{A} \mathbf{x}$:

- f is convex if \mathbf{A} is psd
- f is strictly convex if \mathbf{A} is pd

Subgradients

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

$$f(\mathbf{z}) \geq f(\mathbf{x}) + \mathbf{g}^\top (\mathbf{z} - \mathbf{x}) \quad \forall \mathbf{z} \in \text{dom}(f) \quad (8)$$

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

$$f(\mathbf{z}) \geq f(\mathbf{x}) + \mathbf{g}^\top (\mathbf{z} - \mathbf{x}) \quad \forall \mathbf{z} \in \text{dom}(f) \quad (8)$$

- The set of subgradients is called the **subdifferential**: $\partial f(\mathbf{x})$

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

$$f(\mathbf{z}) \geq f(\mathbf{x}) + \mathbf{g}^\top(\mathbf{z} - \mathbf{x}) \quad \forall \mathbf{z} \in \text{dom}(f) \quad (8)$$

- The set of subgradients is called the **subdifferential**: $\partial f(\mathbf{x})$
- f is subdifferentiable at \mathbf{x} if it has at least one subgradient at that point

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

$$f(\mathbf{z}) \geq f(\mathbf{x}) + \mathbf{g}^\top(\mathbf{z} - \mathbf{x}) \quad \forall \mathbf{z} \in \text{dom}(f) \quad (8)$$

- The set of subgradients is called the **subdifferential**: $\partial f(\mathbf{x})$
- f is subdifferentiable at \mathbf{x} if it has at least one subgradient at that point

Subdifferential of ReLU

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

$$f(\mathbf{z}) \geq f(\mathbf{x}) + \mathbf{g}^\top (\mathbf{z} - \mathbf{x}) \quad \forall \mathbf{z} \in \text{dom}(f) \quad (8)$$

- The set of subgradients is called the **subdifferential**: $\partial f(\mathbf{x})$
- f is subdifferentiable at \mathbf{x} if it has at least one subgradient at that point

Subdifferential of ReLU

The rectified linear unit function (ReLU) is given by:

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

$$f(\mathbf{z}) \geq f(\mathbf{x}) + \mathbf{g}^\top (\mathbf{z} - \mathbf{x}) \quad \forall \mathbf{z} \in \text{dom}(f) \quad (8)$$

- The set of subgradients is called the **subdifferential**: $\partial f(\mathbf{x})$
- f is subdifferentiable at \mathbf{x} if it has at least one subgradient at that point

Subdifferential of ReLU

The rectified linear unit function (ReLU) is given by:

$$\text{ReLU}(z) = \max(0, z), \quad z \in \mathbb{R} \quad (9)$$

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

$$f(\mathbf{z}) \geq f(\mathbf{x}) + \mathbf{g}^\top (\mathbf{z} - \mathbf{x}) \quad \forall \mathbf{z} \in \text{dom}(f) \quad (8)$$

- The set of subgradients is called the **subdifferential**: $\partial f(\mathbf{x})$
- f is subdifferentiable at \mathbf{x} if it has at least one subgradient at that point

Subdifferential of ReLU

The rectified linear unit function (ReLU) is given by:

$$\text{ReLU}(z) = \max(0, z), \quad z \in \mathbb{R} \quad (9)$$

Its subdifferential is:

Subgradients

Given a convex function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, $\mathbf{g} \in \mathbb{R}^n$ is a subgradient of f at $\mathbf{x} \in \text{dom}(f)$ if:

$$f(\mathbf{z}) \geq f(\mathbf{x}) + \mathbf{g}^\top (\mathbf{z} - \mathbf{x}) \quad \forall \mathbf{z} \in \text{dom}(f) \quad (8)$$

- The set of subgradients is called the **subdifferential**: $\partial f(\mathbf{x})$
- f is subdifferentiable at \mathbf{x} if it has at least one subgradient at that point

Subdifferential of ReLU

The rectified linear unit function (ReLU) is given by:

$$\text{ReLU}(z) = \max(0, z), \quad z \in \mathbb{R} \quad (9)$$

Its subdifferential is:

$$\partial \text{ReLU}(z) = \begin{cases} 0 & \text{if } z < 0 \\ [0, 1] & \text{if } z = 0 \\ 1 & \text{if } z > 0 \end{cases} \quad (10)$$

First-order methods

First-order methods

These methods only require first-order derivatives in order to compute an update:

First-order methods

These methods only require first-order derivatives in order to compute an update:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \rho_t \mathbf{d}_t \quad (11)$$

First-order methods

These methods only require first-order derivatives in order to compute an update:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \rho_t \mathbf{d}_t \quad (11)$$

where:

First-order methods

These methods only require first-order derivatives in order to compute an update:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \rho_t \mathbf{d}_t \quad (11)$$

where:

- ρ_t : step size/learning rate

First-order methods

These methods only require first-order derivatives in order to compute an update:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \rho_t \mathbf{d}_t \quad (11)$$

where:

- ρ_t : step size/learning rate
- \mathbf{d}_t : descent direction (e.g. negative gradient)

First-order methods

These methods only require first-order derivatives in order to compute an update:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \rho_t \mathbf{d}_t \quad (11)$$

where:

- ρ_t : step size/learning rate
- \mathbf{d}_t : descent direction (e.g. negative gradient)
- t : index of iteration

First-order methods

These methods only require first-order derivatives in order to compute an update:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \rho_t \mathbf{d}_t \quad (11)$$

where:

- ρ_t : step size/learning rate
- \mathbf{d}_t : descent direction (e.g. negative gradient)
- t : index of iteration

The classic first-order method is **gradient descent**

Descent direction

Descent direction

A vector \mathbf{d} is considered a descent direction if a nonzero step size ρ moved in its direction yields a decrease in the objective:

Descent direction

A vector \mathbf{d} is considered a descent direction if a nonzero step size ρ moved in its direction yields a decrease in the objective:

$$\mathcal{L}(\boldsymbol{\theta} + \rho \mathbf{d}) < \mathcal{L}(\boldsymbol{\theta}) \quad \forall 0 < \rho < \rho_{\max} \quad (12)$$

Descent direction

A vector \mathbf{d} is considered a descent direction if a nonzero step size ρ moved in its direction yields a decrease in the objective:

$$\mathcal{L}(\boldsymbol{\theta} + \rho \mathbf{d}) < \mathcal{L}(\boldsymbol{\theta}) \quad \forall 0 < \rho < \rho_{\max} \quad (12)$$

Steepest descent

Descent direction

A vector \mathbf{d} is considered a descent direction if a nonzero step size ρ moved in its direction yields a decrease in the objective:

$$\mathcal{L}(\boldsymbol{\theta} + \rho \mathbf{d}) < \mathcal{L}(\boldsymbol{\theta}) \quad \forall 0 < \rho < \rho_{\max} \quad (12)$$

Steepest descent

The direction of steepest descent is given by the negative gradient:

Descent direction

A vector \mathbf{d} is considered a descent direction if a nonzero step size ρ moved in its direction yields a decrease in the objective:

$$\mathcal{L}(\boldsymbol{\theta} + \rho \mathbf{d}) < \mathcal{L}(\boldsymbol{\theta}) \quad \forall 0 < \rho < \rho_{\max} \quad (12)$$

Steepest descent

The direction of steepest descent is given by the negative gradient:

$$-\mathbf{g}_t = -\nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (13)$$

Descent direction

A vector \mathbf{d} is considered a descent direction if a nonzero step size ρ moved in its direction yields a decrease in the objective:

$$\mathcal{L}(\boldsymbol{\theta} + \rho \mathbf{d}) < \mathcal{L}(\boldsymbol{\theta}) \quad \forall 0 < \rho < \rho_{\max} \quad (12)$$

Steepest descent

The direction of steepest descent is given by the negative gradient:

$$-\mathbf{g}_t = -\nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (13)$$

Step size

Step size

The choice of step size impacts the convergence of an optimization routine. It determines how far along in the descent direction the variable is updated at each step.

Step size

The choice of step size impacts the convergence of an optimization routine. It determines how far along in the descent direction the variable is updated at each step.

- **Constant step size:** ρ is fixed throughout

Step size

The choice of step size impacts the convergence of an optimization routine. It determines how far along in the descent direction the variable is updated at each step.

- **Constant step size:** ρ is fixed throughout
- **Adaptive step size:** ρ_t is modified at each iteration to satisfy certain conditions (e.g. line search methods)

Step size

The choice of step size impacts the convergence of an optimization routine. It determines how far along in the descent direction the variable is updated at each step.

- **Constant step size:** ρ is fixed throughout
- **Adaptive step size:** ρ_t is modified at each iteration to satisfy certain conditions (e.g. line search methods)
- Sequence $\{\rho_t\}$ in an optimization algorithm is known as the learning rate schedule

Example: Gradient descent

Example: Gradient descent

Given the function

Example: Gradient descent

Given the function

$$f(x) = (x - 1)^4 - 3x + 4$$

Example: Gradient descent

Given the function

$$f(x) = (x - 1)^4 - 3x + 4$$

find the optimal point $(x^*, f(x^*))$.

Example: Gradient descent

Given the function

$$f(x) = (x - 1)^4 - 3x + 4$$

find the optimal point $(x^*, f(x^*))$.

First, we find the **gradient** and define the **update** step:

Example: Gradient descent

Given the function

$$f(x) = (x - 1)^4 - 3x + 4$$

find the optimal point $(x^*, f(x^*))$.

First, we find the **gradient** and define the **update** step:

$$f'(x) =$$

Example: Gradient descent

Given the function

$$f(x) = (x - 1)^4 - 3x + 4$$

find the optimal point $(x^*, f(x^*))$.

First, we find the **gradient** and define the **update** step:

$$f'(x) = 4(x - 1)^3 - 3$$

Example: Gradient descent

Given the function

$$f(x) = (x - 1)^4 - 3x + 4$$

find the optimal point $(x^*, f(x^*))$.

First, we find the **gradient** and define the **update** step:

$$f'(x) = 4(x - 1)^3 - 3$$

$$x_{k+1} =$$

Example: Gradient descent

Given the function

$$f(x) = (x - 1)^4 - 3x + 4$$

find the optimal point $(x^*, f(x^*))$.


First, we find the **gradient** and define the **update** step:

$$\begin{aligned} f'(x) &= 4(x - 1)^3 - 3 \\ x_{k+1} &= x_k - \lambda [4(x_k - 1)^3 - 3] \end{aligned}$$

We choose $\lambda = 0.05$ and a random starting point between 0 and 1.


Example: Gradient descent (cont.)

Example: Gradient descent (cont.)



Using the `gradient-descent.ipynb` notebook, we find the optimal point as **(1.908, -1.044)**.

Example: Gradient descent (cont.)



Using the `gradient-descent.ipynb` notebook, we find the optimal point as **(1.908, -1.044)**. How does the learning rate impact the solution process?

Exact line search

Exact line search

Finds optimal ρ_t by:

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* =$$

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho)$$

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho) = \arg \min_{\rho > 0} \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \quad (14)$$

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho) = \arg \min_{\rho > 0} \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \quad (14)$$

If $\mathcal{L}(\boldsymbol{\theta})$ is quadratic, then we can write:

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho) = \arg \min_{\rho > 0} \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \quad (14)$$

If $\mathcal{L}(\boldsymbol{\theta})$ is quadratic, then we can write:

$$\phi_t(\rho)$$

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho) = \arg \min_{\rho > 0} \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \quad (14)$$

If $\mathcal{L}(\boldsymbol{\theta})$ is quadratic, then we can write:

$$\phi_t(\rho) = \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t)$$

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho) = \arg \min_{\rho > 0} \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \quad (14)$$

If $\mathcal{L}(\boldsymbol{\theta})$ is quadratic, then we can write:

$$\phi_t(\rho) = \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) = \frac{1}{2}(\boldsymbol{\theta} + \rho \mathbf{d})^\top \mathbf{A}(\boldsymbol{\theta} + \rho \mathbf{d}) + \mathbf{b}^\top(\boldsymbol{\theta} + \rho \mathbf{d}) + c \quad (15)$$

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho) = \arg \min_{\rho > 0} \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \quad (14)$$

If $\mathcal{L}(\boldsymbol{\theta})$ is quadratic, then we can write:

$$\phi_t(\rho) = \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) = \frac{1}{2}(\boldsymbol{\theta} + \rho \mathbf{d})^\top \mathbf{A}(\boldsymbol{\theta} + \rho \mathbf{d}) + \mathbf{b}^\top(\boldsymbol{\theta} + \rho \mathbf{d}) + c \quad (15)$$

Solving for $\phi'_t(\rho) = 0$ gives the optimal update as:

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho) = \arg \min_{\rho > 0} \mathcal{L}(\theta_t + \rho \mathbf{d}_t) \quad (14)$$

If $\mathcal{L}(\theta)$ is quadratic, then we can write:

$$\phi_t(\rho) = \mathcal{L}(\theta_t + \rho \mathbf{d}_t) = \frac{1}{2}(\theta + \rho \mathbf{d})^\top \mathbf{A}(\theta + \rho \mathbf{d}) + \mathbf{b}^\top(\theta + \rho \mathbf{d}) + c \quad (15)$$

Solving for $\phi'_t(\rho) = 0$ gives the optimal update as:

$$\rho_t^* = -\frac{\mathbf{d}_t^\top (\mathbf{A}\theta_t + \mathbf{b})}{\mathbf{d}_t^\top \mathbf{A} \mathbf{d}_t} \quad (16)$$

Exact line search

Finds optimal ρ_t by:

$$\rho_t^* = \arg \min_{\rho > 0} \phi_t(\rho) = \arg \min_{\rho > 0} \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \quad (14)$$

If $\mathcal{L}(\boldsymbol{\theta})$ is quadratic, then we can write:

$$\phi_t(\rho) = \mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) = \frac{1}{2}(\boldsymbol{\theta} + \rho \mathbf{d})^\top \mathbf{A}(\boldsymbol{\theta} + \rho \mathbf{d}) + \mathbf{b}^\top(\boldsymbol{\theta} + \rho \mathbf{d}) + c \quad (15)$$

Solving for $\phi'_t(\rho) = 0$ gives the optimal update as:

$$\rho_t^* = -\frac{\mathbf{d}_t^\top (\mathbf{A}\boldsymbol{\theta}_t + \mathbf{b})}{\mathbf{d}_t^\top \mathbf{A} \mathbf{d}_t} \quad (16)$$

This approach introduces additional computational expense

Armijo backtracking

Finds a suitable ρ_t^* by iteratively shrinking ρ_t by $\beta \in (0, 1)$, i.e.

Armijo backtracking

Finds a suitable ρ_t^* by iteratively shrinking ρ_t by $\beta \in (0, 1)$, i.e.

$$\rho_t^{k+1} = \beta \rho_t^k \quad (17)$$

until the **Armijo-Goldstein** condition is satisfied:

Armijo backtracking

Finds a suitable ρ_t^* by iteratively shrinking ρ_t by $\beta \in (0, 1)$, i.e.

$$\rho_t^{k+1} = \beta \rho_t^k \quad (17)$$

until the **Armijo-Goldstein** condition is satisfied:

$$\mathcal{L}(\theta_t + \rho \mathbf{d}_t) \leq \mathcal{L}(\theta_t) + c \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\theta_t) \quad (18)$$

Armijo backtracking

Finds a suitable ρ_t^* by iteratively shrinking ρ_t by $\beta \in (0, 1)$, i.e.

$$\rho_t^{k+1} = \beta \rho_t^k \quad (17)$$

until the **Armijo-Goldstein** condition is satisfied:

$$\mathcal{L}(\theta_t + \rho \mathbf{d}_t) \leq \mathcal{L}(\theta_t) + c \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\theta_t) \quad (18)$$

where $c \in [0, 1]$ is typically chosen as 10^{-4}

Momentum

Momentum

Momentum methods are used to improve convergence.

Momentum

Momentum methods are used to improve convergence. The momentum update is given by:

Momentum

Momentum methods are used to improve convergence. The momentum update is given by:

$$\mathbf{m}_{t+1} =$$

Momentum

Momentum methods are used to improve convergence. The momentum update is given by:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t + \mathbf{g}_t \quad (19)$$

Momentum

Momentum methods are used to improve convergence. The momentum update is given by:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t + \mathbf{g}_t \quad (19)$$

where $\beta \in (0, 1)$.

Momentum

Momentum methods are used to improve convergence. The momentum update is given by:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t + \mathbf{g}_t \quad (19)$$

where $\beta \in (0, 1)$.

The momentum can also be expressed as an exponentially weighted average of past gradients:

Momentum

Momentum methods are used to improve convergence. The momentum update is given by:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t + \mathbf{g}_t \quad (19)$$

where $\beta \in (0, 1)$.

The momentum can also be expressed as an exponentially weighted average of past gradients:

$$\mathbf{m}_{t+1} = \sum_{\tau=0}^t \beta^\tau \mathbf{g}_{t-\tau} \quad (20)$$

Momentum

Momentum methods are used to improve convergence. The momentum update is given by:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t + \mathbf{g}_t \quad (19)$$

where $\beta \in (0, 1)$.

The momentum can also be expressed as an exponentially weighted average of past gradients:

$$\mathbf{m}_{t+1} = \sum_{\tau=0}^t \beta^\tau \mathbf{g}_{t-\tau} \quad (20)$$

The update is given by:

Momentum

Momentum methods are used to improve convergence. The momentum update is given by:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t + \mathbf{g}_t \quad (19)$$

where $\beta \in (0, 1)$.

The momentum can also be expressed as an exponentially weighted average of past gradients:

$$\mathbf{m}_{t+1} = \sum_{\tau=0}^t \beta^\tau \mathbf{g}_{t-\tau} \quad (20)$$

The update is given by:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \rho_t \mathbf{m}_{t+1} \quad (21)$$

Nesterov momentum

Nesterov momentum

In the Nesterov method, the momentum update computes the gradient at the new location, which can speed up convergence:

$$\mathbf{m}_{t+1} =$$

Nesterov momentum

In the Nesterov method, the momentum update computes the gradient at the new location, which can speed up convergence:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t - \rho_t \nabla \mathcal{L}(\boldsymbol{\theta}_t + \beta \mathbf{m}_t) \quad (22)$$

Nesterov momentum

In the Nesterov method, the momentum update computes the gradient at the new location, which can speed up convergence:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t - \rho_t \nabla \mathcal{L}(\boldsymbol{\theta}_t + \beta \mathbf{m}_t) \quad (22)$$

The update is given by:

Nesterov momentum

In the Nesterov method, the momentum update computes the gradient at the new location, which can speed up convergence:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t - \rho_t \nabla \mathcal{L}(\boldsymbol{\theta}_t + \beta \mathbf{m}_t) \quad (22)$$

The update is given by:

$$\boldsymbol{\theta}_{t+1}$$

Nesterov momentum

In the Nesterov method, the momentum update computes the gradient at the new location, which can speed up convergence:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t - \rho_t \nabla \mathcal{L}(\theta_t + \beta \mathbf{m}_t) \quad (22)$$

The update is given by:

$$\theta_{t+1} = \theta_t + \mathbf{m}_{t+1} \quad (23)$$

Nesterov momentum

In the Nesterov method, the momentum update computes the gradient at the new location, which can speed up convergence:

$$\mathbf{m}_{t+1} = \beta \mathbf{m}_t - \rho_t \nabla \mathcal{L}(\boldsymbol{\theta}_t + \beta \mathbf{m}_t) \quad (22)$$

The update is given by:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \mathbf{m}_{t+1} \quad (23)$$

This method is also called **Nesterov accelerated gradient**

Stochastic gradient descent (SGD)

In stochastic optimization, goal is to minimize:

Stochastic gradient descent (SGD)

In stochastic optimization, goal is to minimize:

$$\mathcal{L}(\boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\mathcal{L}(\boldsymbol{\theta}, \mathbf{z})] \quad (24)$$

Stochastic gradient descent (SGD)

In stochastic optimization, goal is to minimize:

$$\mathcal{L}(\boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\mathcal{L}(\boldsymbol{\theta}, \mathbf{z})] \quad (24)$$

where $\mathbf{z}_t \sim q$ could be a training sample drawn from a set.

Stochastic gradient descent (SGD)

In stochastic optimization, goal is to minimize:

$$\mathcal{L}(\boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\mathcal{L}(\boldsymbol{\theta}, \mathbf{z})] \quad (24)$$

where $\mathbf{z}_t \sim q$ could be a training sample drawn from a set. Thus, the **stochastic gradient descent** update is given by:

Stochastic gradient descent (SGD)

In stochastic optimization, goal is to minimize:

$$\mathcal{L}(\boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\mathcal{L}(\boldsymbol{\theta}, \mathbf{z})] \quad (24)$$

where $\mathbf{z}_t \sim q$ could be a training sample drawn from a set. Thus, the **stochastic gradient descent** update is given by:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \rho_t \nabla \mathcal{L}(\boldsymbol{\theta}_t, \mathbf{z}_t) =$$

Stochastic gradient descent (SGD)

In stochastic optimization, goal is to minimize:

$$\mathcal{L}(\boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\mathcal{L}(\boldsymbol{\theta}, \mathbf{z})] \quad (24)$$

where $\mathbf{z}_t \sim q$ could be a training sample drawn from a set. Thus, the **stochastic gradient descent** update is given by:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \rho_t \nabla \mathcal{L}(\boldsymbol{\theta}_t, \mathbf{z}_t) = \boldsymbol{\theta}_t - \rho_t \mathbf{g}_t \quad (25)$$

Stochastic gradient descent (SGD)

In stochastic optimization, goal is to minimize:

$$\mathcal{L}(\boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\mathcal{L}(\boldsymbol{\theta}, \mathbf{z})] \quad (24)$$

where $\mathbf{z}_t \sim q$ could be a training sample drawn from a set. Thus, the **stochastic gradient descent** update is given by:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \rho_t \nabla \mathcal{L}(\boldsymbol{\theta}_t, \mathbf{z}_t) = \boldsymbol{\theta}_t - \rho_t \mathbf{g}_t \quad (25)$$

- Finite sum: $\mathcal{L}(\boldsymbol{\theta}_t) = \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\boldsymbol{\theta}_t)$

Newton's method

Newton's method

Second-order methods incorporate curvature via the second derivative to speed up convergence.

Newton's method

Second-order methods incorporate curvature via the second derivative to speed up convergence.

The basic method of this form is Newton's method:

Newton's method

Second-order methods incorporate curvature via the second derivative to speed up convergence.

The basic method of this form is Newton's method:

$$\theta_{t+1}$$

Newton's method

Second-order methods incorporate curvature via the second derivative to speed up convergence.

The basic method of this form is Newton's method:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \rho_t \mathbf{H}^{-1} \mathbf{g}_t \quad (26)$$

Newton's method

Second-order methods incorporate curvature via the second derivative to speed up convergence.

The basic method of this form is Newton's method:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \rho_t \mathbf{H}^{-1} \mathbf{g}_t \quad (26)$$

where:

Newton's method

Second-order methods incorporate curvature via the second derivative to speed up convergence.

The basic method of this form is Newton's method:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \rho_t \mathbf{H}^{-1} \mathbf{g}_t \quad (26)$$

where:

$$\mathbf{H}_t := \nabla^2 \mathcal{L}(\boldsymbol{\theta}_t)$$

Newton's method

Second-order methods incorporate curvature via the second derivative to speed up convergence.

The basic method of this form is Newton's method:

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \rho_t \mathbf{H}^{-1} \mathbf{g}_t \quad (26)$$

where:

$$\mathbf{H}_t := \nabla^2 \mathcal{L}(\boldsymbol{\theta}_t) = \mathbf{H}(\boldsymbol{\theta}_t) \quad (27)$$

Derivation of Newton's method

Derivation of Newton's method

Taylor-expand $\mathcal{L}(\boldsymbol{\theta})$ around $\boldsymbol{\theta}_t$:

Derivation of Newton's method

Taylor-expand $\mathcal{L}(\boldsymbol{\theta})$ around $\boldsymbol{\theta}_t$:

$$\mathcal{L}_{\text{quad}} = \mathcal{L}(\boldsymbol{\theta}_t) + \mathbf{g}_t^\top (\boldsymbol{\theta} - \boldsymbol{\theta}_t) + \frac{1}{2}(\boldsymbol{\theta} - \boldsymbol{\theta}_t)^\top \mathbf{H}_t(\boldsymbol{\theta} - \boldsymbol{\theta}_t) \quad (28)$$

Derivation of Newton's method

Taylor-expand $\mathcal{L}(\boldsymbol{\theta})$ around $\boldsymbol{\theta}_t$:

$$\mathcal{L}_{\text{quad}} = \mathcal{L}(\boldsymbol{\theta}_t) + \mathbf{g}_t^\top (\boldsymbol{\theta} - \boldsymbol{\theta}_t) + \frac{1}{2}(\boldsymbol{\theta} - \boldsymbol{\theta}_t)^\top \mathbf{H}_t (\boldsymbol{\theta} - \boldsymbol{\theta}_t) \quad (28)$$

Taking the derivative of $\mathcal{L}_{\text{quad}}$ and setting it to zero to find the minimum, we obtain:

Derivation of Newton's method

Taylor-expand $\mathcal{L}(\boldsymbol{\theta})$ around $\boldsymbol{\theta}_t$:

$$\mathcal{L}_{\text{quad}} = \mathcal{L}(\boldsymbol{\theta}_t) + \mathbf{g}_t^\top (\boldsymbol{\theta} - \boldsymbol{\theta}_t) + \frac{1}{2}(\boldsymbol{\theta} - \boldsymbol{\theta}_t)^\top \mathbf{H}_t(\boldsymbol{\theta} - \boldsymbol{\theta}_t) \quad (28)$$

Taking the derivative of $\mathcal{L}_{\text{quad}}$ and setting it to zero to find the minimum, we obtain:

$$\mathbf{g}_t + \mathbf{H}_t(\boldsymbol{\theta} - \boldsymbol{\theta}_t) = \mathbf{0} \quad (29)$$

$$\boldsymbol{\theta} - \boldsymbol{\theta}_t = -\mathbf{H}_t^{-1} \mathbf{g}_t \quad (30)$$

$$\boldsymbol{\theta} = \boldsymbol{\theta}_t - \mathbf{H}_t^{-1} \mathbf{g}_t \quad (31)$$

Derivation of Newton's method

Taylor-expand $\mathcal{L}(\boldsymbol{\theta})$ around $\boldsymbol{\theta}_t$:

$$\mathcal{L}_{\text{quad}} = \mathcal{L}(\boldsymbol{\theta}_t) + \mathbf{g}_t^\top (\boldsymbol{\theta} - \boldsymbol{\theta}_t) + \frac{1}{2}(\boldsymbol{\theta} - \boldsymbol{\theta}_t)^\top \mathbf{H}_t(\boldsymbol{\theta} - \boldsymbol{\theta}_t) \quad (28)$$

Taking the derivative of $\mathcal{L}_{\text{quad}}$ and setting it to zero to find the minimum, we obtain:

$$\mathbf{g}_t + \mathbf{H}_t(\boldsymbol{\theta} - \boldsymbol{\theta}_t) = \mathbf{0} \quad (29)$$

$$\boldsymbol{\theta} - \boldsymbol{\theta}_t = -\mathbf{H}_t^{-1} \mathbf{g}_t \quad (30)$$

$$\boldsymbol{\theta} = \boldsymbol{\theta}_t - \mathbf{H}_t^{-1} \mathbf{g}_t \quad (31)$$

Thus, we set the descent direction \mathbf{d}_t as $-\mathbf{H}_t^{-1} \mathbf{g}_t$

BFGS

BFGS

Broyden-Fletcher-Goldfarb-Shanno method.

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

$$\mathbf{B}_{t+1} = \mathbf{B}_t + \frac{\mathbf{y}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} - \frac{(\mathbf{B}_t \mathbf{s}_t)(\mathbf{B}_t \mathbf{s}_t)^\top}{\mathbf{s}_t^\top \mathbf{B}_t \mathbf{s}_t} \quad (32)$$

$$\mathbf{s}_t = \boldsymbol{\theta}_t - \boldsymbol{\theta}_{t-1} \quad (33)$$

$$\mathbf{y}_t = \mathbf{g}_t - \mathbf{g}_{t-1} \quad (34)$$

- \mathbf{B}_0 is typically initialized as \mathbf{I} (positive definite)
- To ensure \mathbf{B}_{t+1} remains pd, ρ must satisfy the Wolfe conditions:

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \leq \mathcal{L}(\boldsymbol{\theta}_t) + c_1 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (35)$$

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \geq c_2 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (36)$$

where $0 < c_1 < c_2 < 1$

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

$$\mathbf{B}_{t+1} = \mathbf{B}_t + \frac{\mathbf{y}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} - \frac{(\mathbf{B}_t \mathbf{s}_t)(\mathbf{B}_t \mathbf{s}_t)^\top}{\mathbf{s}_t^\top \mathbf{B}_t \mathbf{s}_t} \quad (32)$$

$$\mathbf{y}_t = \mathbf{g}_t - \mathbf{g}_{t-1} \quad (34)$$

- \mathbf{B}_0 is typically initialized as \mathbf{I} (positive definite)
- To ensure \mathbf{B}_{t+1} remains pd, ρ must satisfy the Wolfe conditions:

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \leq \mathcal{L}(\boldsymbol{\theta}_t) + c_1 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (35)$$

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \geq c_2 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (36)$$

where $0 < c_1 < c_2 < 1$

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

$$\mathbf{B}_{t+1} = \mathbf{B}_t + \frac{\mathbf{y}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} - \frac{(\mathbf{B}_t \mathbf{s}_t)(\mathbf{B}_t \mathbf{s}_t)^\top}{\mathbf{s}_t^\top \mathbf{B}_t \mathbf{s}_t} \quad (32)$$

$$\mathbf{s}_t = \boldsymbol{\theta}_t - \boldsymbol{\theta}_{t-1} \quad (33)$$

$$\mathbf{y}_t = \mathbf{g}_t - \mathbf{g}_{t-1} \quad (34)$$

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

$$\mathbf{B}_{t+1} = \mathbf{B}_t + \frac{\mathbf{y}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} - \frac{(\mathbf{B}_t \mathbf{s}_t)(\mathbf{B}_t \mathbf{s}_t)^\top}{\mathbf{s}_t^\top \mathbf{B}_t \mathbf{s}_t} \quad (32)$$

$$\mathbf{s}_t = \boldsymbol{\theta}_t - \boldsymbol{\theta}_{t-1} \quad (33)$$

$$\mathbf{y}_t = \mathbf{g}_t - \mathbf{g}_{t-1} \quad (34)$$

- \mathbf{B}_0 is typically initialized as \mathbf{I} (positive definite)

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

$$\mathbf{B}_{t+1} = \mathbf{B}_t + \frac{\mathbf{y}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} - \frac{(\mathbf{B}_t \mathbf{s}_t)(\mathbf{B}_t \mathbf{s}_t)^\top}{\mathbf{s}_t^\top \mathbf{B}_t \mathbf{s}_t} \quad (32)$$

$$\mathbf{s}_t = \boldsymbol{\theta}_t - \boldsymbol{\theta}_{t-1} \quad (33)$$

$$\mathbf{y}_t = \mathbf{g}_t - \mathbf{g}_{t-1} \quad (34)$$

- \mathbf{B}_0 is typically initialized as \mathbf{I} (positive definite)
- To ensure \mathbf{B}_{t+1} remains pd, ρ must satisfy the Wolfe conditions:

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

$$\mathbf{B}_{t+1} = \mathbf{B}_t + \frac{\mathbf{y}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} - \frac{(\mathbf{B}_t \mathbf{s}_t)(\mathbf{B}_t \mathbf{s}_t)^\top}{\mathbf{s}_t^\top \mathbf{B}_t \mathbf{s}_t} \quad (32)$$

$$\mathbf{s}_t = \boldsymbol{\theta}_t - \boldsymbol{\theta}_{t-1} \quad (33)$$

$$\mathbf{y}_t = \mathbf{g}_t - \mathbf{g}_{t-1} \quad (34)$$

- \mathbf{B}_0 is typically initialized as \mathbf{I} (positive definite)
- To ensure \mathbf{B}_{t+1} remains pd, ρ must satisfy the Wolfe conditions:

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \leq \mathcal{L}(\boldsymbol{\theta}_t) + c_1 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (35)$$

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \geq c_2 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (36)$$

where $0 < c_1 < c_2 < 1$

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

$$\mathbf{B}_{t+1} = \mathbf{B}_t + \frac{\mathbf{y}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} - \frac{(\mathbf{B}_t \mathbf{s}_t)(\mathbf{B}_t \mathbf{s}_t)^\top}{\mathbf{s}_t^\top \mathbf{B}_t \mathbf{s}_t} \quad (32)$$

$$\mathbf{s}_t = \boldsymbol{\theta}_t - \boldsymbol{\theta}_{t-1} \quad (33)$$

$$\mathbf{y}_t = \mathbf{g}_t - \mathbf{g}_{t-1} \quad (34)$$

- \mathbf{B}_0 is typically initialized as \mathbf{I} (positive definite)
- To ensure \mathbf{B}_{t+1} remains pd, ρ must satisfy the Wolfe conditions:

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \leq \mathcal{L}(\boldsymbol{\theta}_t) + c_1 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (35)$$

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \geq c_2 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (36)$$

BFGS

Broyden-Fletcher-Goldfarb-Shanno method. Approximates $\mathbf{H}_t \approx \mathbf{B}_t$:

$$\mathbf{B}_{t+1} = \mathbf{B}_t + \frac{\mathbf{y}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} - \frac{(\mathbf{B}_t \mathbf{s}_t)(\mathbf{B}_t \mathbf{s}_t)^\top}{\mathbf{s}_t^\top \mathbf{B}_t \mathbf{s}_t} \quad (32)$$

$$\mathbf{s}_t = \boldsymbol{\theta}_t - \boldsymbol{\theta}_{t-1} \quad (33)$$

$$\mathbf{y}_t = \mathbf{g}_t - \mathbf{g}_{t-1} \quad (34)$$

- \mathbf{B}_0 is typically initialized as \mathbf{I} (positive definite)
- To ensure \mathbf{B}_{t+1} remains pd, ρ must satisfy the Wolfe conditions:

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \leq \mathcal{L}(\boldsymbol{\theta}_t) + c_1 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (35)$$

$$\mathcal{L}(\boldsymbol{\theta}_t + \rho \mathbf{d}_t) \geq c_2 \rho \mathbf{d}_t^\top \nabla \mathcal{L}(\boldsymbol{\theta}_t) \quad (36)$$

where $0 < c_1 < c_2 < 1$

Limited memory BFGS

Limited memory BFGS

In practice, BFGS directly provides the inverse Hessian approximation:

Limited memory BFGS

In practice, BFGS directly provides the inverse Hessian approximation:

$$\mathbf{C}_{t+1} = \left(\mathbf{I} - \frac{\mathbf{s}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \right) \mathbf{C}_t \left(\mathbf{I} - \frac{\mathbf{y}_t \mathbf{s}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \right) + \frac{\mathbf{s}_t \mathbf{s}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \quad (37)$$

Limited memory BFGS

In practice, BFGS directly provides the inverse Hessian approximation:

$$\mathbf{C}_{t+1} = \left(\mathbf{I} - \frac{\mathbf{s}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \right) \mathbf{C}_t \left(\mathbf{I} - \frac{\mathbf{y}_t \mathbf{s}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \right) + \frac{\mathbf{s}_t \mathbf{s}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \quad (37)$$

where $\mathbf{C}_t \approx \mathbf{H}_t^{-1}$.

Limited memory BFGS

In practice, BFGS directly provides the inverse Hessian approximation:

$$\mathbf{C}_{t+1} = \left(\mathbf{I} - \frac{\mathbf{s}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \right) \mathbf{C}_t \left(\mathbf{I} - \frac{\mathbf{y}_t \mathbf{s}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \right) + \frac{\mathbf{s}_t \mathbf{s}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \quad (37)$$

where $\mathbf{C}_t \approx \mathbf{H}_t^{-1}$.

- The update can be reduced to a recurrence relation that depends explicitly on \mathbf{C}_0 and the history of \mathbf{s}_t and \mathbf{y}_t

Limited memory BFGS

In practice, BFGS directly provides the inverse Hessian approximation:

$$\mathbf{C}_{t+1} = \left(\mathbf{I} - \frac{\mathbf{s}_t \mathbf{y}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \right) \mathbf{C}_t \left(\mathbf{I} - \frac{\mathbf{y}_t \mathbf{s}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \right) + \frac{\mathbf{s}_t \mathbf{s}_t^\top}{\mathbf{y}_t^\top \mathbf{s}_t} \quad (37)$$

where $\mathbf{C}_t \approx \mathbf{H}_t^{-1}$.

- The update can be reduced to a recurrence relation that depends explicitly on \mathbf{C}_0 and the history of \mathbf{s}_t and \mathbf{y}_t
- For computational efficiency, the M most recent $\mathbf{s}_t, \mathbf{y}_t$ may be used instead
- M is typically $\in [5, 20]$
- This approach is termed L-BFGS (limited memory BFGS)

Maximum likelihood estimation: Key definitions

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta)$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta)$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta) \cdots$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

The **maximum likelihood estimator (MLE)**, $\hat{\theta}$,

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

The **maximum likelihood estimator (MLE)**, $\hat{\theta}$, is the value of θ that maximizes the [log-]likelihood function:

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta) f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

The **maximum likelihood estimator (MLE)**, $\hat{\theta}$, is the value of θ that maximizes the [log-]likelihood function:

$$\frac{\partial \ln L(x_1, x_2, \dots, x_n; \theta)}{\partial \theta} =$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

The **maximum likelihood estimator (MLE)**, $\hat{\theta}$, is the value of θ that maximizes the [log-]likelihood function:

$$\frac{\partial \ln L(x_1, x_2, \dots, x_n; \theta)}{\partial \theta} = 0 \quad (39)$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

The **maximum likelihood estimator (MLE)**, $\hat{\theta}$, is the value of θ that maximizes the [log-]likelihood function:

$$\frac{\partial \ln L(x_1, x_2, \dots, x_n; \theta)}{\partial \theta} = 0 \quad (39)$$

For multiple parameters, the likelihood function is:

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

The **maximum likelihood estimator (MLE)**, $\hat{\theta}$, is the value of θ that maximizes the [log-]likelihood function:

$$\frac{\partial \ln L(x_1, x_2, \dots, x_n; \theta)}{\partial \theta} = 0 \quad (39)$$

For multiple parameters, the likelihood function is:

$$L(x_1, \dots, x_n; \theta_1, \dots, \theta_1, \dots, \theta_m)$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta) f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

The **maximum likelihood estimator (MLE)**, $\hat{\theta}$, is the value of θ that maximizes the [log-]likelihood function:

$$\frac{\partial \ln L(x_1, x_2, \dots, x_n; \theta)}{\partial \theta} = 0 \quad (39)$$

For multiple parameters, the likelihood function is:

$$L(x_1, \dots, x_n; \theta_1, \dots, \theta_m) = \prod_{i=1}^n f(x_i; \theta_i; \theta_1, \dots, \theta_m) \quad (40)$$

Maximum likelihood estimation: Key definitions

Given a set of n independent observations x_1, x_2, \dots, x_n from a random sample, the **likelihood function** is:

$$L(x_1, x_2, \dots, x_n; \theta) = f(x_1; \theta)f(x_2; \theta) \cdots f(x_n; \theta) \quad (38)$$

The **maximum likelihood estimator (MLE)**, $\hat{\theta}$, is the value of θ that maximizes the [log-]likelihood function:

$$\frac{\partial \ln L(x_1, x_2, \dots, x_n; \theta)}{\partial \theta} = 0 \quad (39)$$

For multiple parameters, the likelihood function is:

$$L(x_1, \dots, x_n; \theta_1, \dots, \theta_m) = \prod_{i=1}^n f(x_i; \theta_i; \theta_1, \dots, \theta_m) \quad (40)$$

And the MLE's would be found by simultaneously solving the partial derivatives set to 0 for each parameter.

Maximum likelihood estimation (MLE) in logistic regression

Maximum likelihood estimation (MLE) in logistic regression

The log-likelihood function for the binomial logistic regression case is

Maximum likelihood estimation (MLE) in logistic regression

The log-likelihood function for the binomial logistic regression case is

$$\ell(\beta) = \sum_i [y_i (\beta_0 + \beta_1 x_i) - \log (1 + e^{\beta_0 + \beta_1 x_i})] \quad (41)$$

Maximum likelihood estimation (MLE) in logistic regression

The log-likelihood function for the binomial logistic regression case is

$$\ell(\beta) = \sum_i [y_i (\beta_0 + \beta_1 x_i) - \log (1 + e^{\beta_0 + \beta_1 x_i})] \quad (41)$$

The optimal $\hat{\beta}$ which maximizes $\ell(\beta)$ is the **maximum likelihood estimate**.

Maximum likelihood estimation (MLE) in logistic regression

The log-likelihood function for the binomial logistic regression case is

$$\ell(\beta) = \sum_i [y_i (\beta_0 + \beta_1 x_i) - \log (1 + e^{\beta_0 + \beta_1 x_i})] \quad (41)$$

The optimal $\hat{\beta}$ which maximizes $\ell(\beta)$ is the **maximum likelihood estimate**.

Also recall the derivative of ℓ :

Maximum likelihood estimation (MLE) in logistic regression

The log-likelihood function for the binomial logistic regression case is

$$\ell(\beta) = \sum_i [y_i (\beta_0 + \beta_1 x_i) - \log (1 + e^{\beta_0 + \beta_1 x_i})] \quad (41)$$

The optimal $\hat{\beta}$ which maximizes $\ell(\beta)$ is the **maximum likelihood estimate**.

Also recall the derivative of ℓ :

$$\nabla_{\beta} \ell =$$

Maximum likelihood estimation (MLE) in logistic regression

The log-likelihood function for the binomial logistic regression case is

$$\ell(\beta) = \sum_i [y_i (\beta_0 + \beta_1 x_i) - \log (1 + e^{\beta_0 + \beta_1 x_i})] \quad (41)$$

The optimal $\hat{\beta}$ which maximizes $\ell(\beta)$ is the **maximum likelihood estimate**.

Also recall the derivative of ℓ :

$$\nabla_{\beta} \ell = \begin{pmatrix} \frac{\partial \ell}{\partial \beta_0} \\ \frac{\partial \ell}{\partial \beta_1} \end{pmatrix} =$$

Maximum likelihood estimation (MLE) in logistic regression

The log-likelihood function for the binomial logistic regression case is

$$\ell(\beta) = \sum_i [y_i (\beta_0 + \beta_1 x_i) - \log (1 + e^{\beta_0 + \beta_1 x_i})] \quad (41)$$

The optimal $\hat{\beta}$ which maximizes $\ell(\beta)$ is the **maximum likelihood estimate**.

Also recall the derivative of ℓ :

$$\nabla_{\beta} \ell = \begin{pmatrix} \frac{\partial \ell}{\partial \beta_0} \\ \frac{\partial \ell}{\partial \beta_1} \end{pmatrix} = \begin{pmatrix} \sum_i [y_i - p(x_i)] \\ \sum_i [x_i (y_i - p(x_i))] \end{pmatrix} \quad (42)$$

Maximum likelihood estimation (MLE) in logistic regression

The log-likelihood function for the binomial logistic regression case is

$$\ell(\beta) = \sum_i [y_i (\beta_0 + \beta_1 x_i) - \log (1 + e^{\beta_0 + \beta_1 x_i})] \quad (41)$$

The optimal $\hat{\beta}$ which maximizes $\ell(\beta)$ is the **maximum likelihood estimate**.

Also recall the derivative of ℓ :

$$\nabla_{\beta} \ell = \begin{pmatrix} \frac{\partial \ell}{\partial \beta_0} \\ \frac{\partial \ell}{\partial \beta_1} \end{pmatrix} = \begin{pmatrix} \sum_i [y_i - p(x_i)] \\ \sum_i [x_i (y_i - p(x_i))] \end{pmatrix} \quad (42)$$

We can use either Newton-Raphson or gradient *ascent* to *maximize* ℓ .

Gradient ascent for MLE in logistic regression

Gradient ascent for MLE in logistic regression

This approach only requires the first derivative:

Gradient ascent for MLE in logistic regression

This approach only requires the first derivative:

$$\beta_{k+1} = \beta_k + \lambda \nabla \ell(\beta_k) \quad (43)$$

Gradient ascent for MLE in logistic regression

This approach only requires the first derivative:

$$\beta_{k+1} = \beta_k + \lambda \nabla \ell(\beta_k) \quad (43)$$

Thus, to find $\hat{\beta}$ we iterate using:

Gradient ascent for MLE in logistic regression

This approach only requires the first derivative:

$$\beta_{k+1} = \beta_k + \lambda \nabla \ell(\beta_k) \quad (43)$$

Thus, to find $\hat{\beta}$ we iterate using:

$$\begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} =$$

Gradient ascent for MLE in logistic regression

This approach only requires the first derivative:

$$\beta_{k+1} = \beta_k + \lambda \nabla \ell(\beta_k) \quad (43)$$

Thus, to find $\hat{\beta}$ we iterate using:

$$\begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} = \begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} + \lambda \begin{pmatrix} \sum_i [y_i - p(x_i)] \\ \sum_i [x_i (y_i - p(x_i))] \end{pmatrix} \quad (44)$$

Gradient ascent for MLE in logistic regression

This approach only requires the first derivative:

$$\beta_{k+1} = \beta_k + \lambda \nabla \ell(\beta_k) \quad (43)$$

Thus, to find $\hat{\beta}$ we iterate using:

$$\begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} = \begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} + \lambda \begin{pmatrix} \sum_i [y_i - p(x_i)] \\ \sum_i [x_i (y_i - p(x_i))] \end{pmatrix} \quad (44)$$

Because the log-likelihood is *concave*, and thus a *maximization* problem, we *ascend* the function and thus *add* the scaled derivative.

Gradient ascent for MLE in logistic regression

This approach only requires the first derivative:

$$\beta_{k+1} = \beta_k + \lambda \nabla \ell(\beta_k) \quad (43)$$

Thus, to find $\hat{\beta}$ we iterate using:

$$\begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} = \begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} + \lambda \begin{pmatrix} \sum_i [y_i - p(x_i)] \\ \sum_i [x_i (y_i - p(x_i))] \end{pmatrix} \quad (44)$$

Because the log-likelihood is *concave*, and thus a *maximization* problem, we *ascend* the function and thus *add* the scaled derivative.

- As we can see, the gradient ascent method does not require a second derivative

Gradient ascent for MLE in logistic regression

This approach only requires the first derivative:

$$\beta_{k+1} = \beta_k + \lambda \nabla \ell(\beta_k) \quad (43)$$

Thus, to find $\hat{\beta}$ we iterate using:

$$\begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} = \begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} + \lambda \begin{pmatrix} \sum_i [y_i - p(x_i)] \\ \sum_i [x_i (y_i - p(x_i))] \end{pmatrix} \quad (44)$$

Because the log-likelihood is *concave*, and thus a *maximization* problem, we *ascend* the function and thus *add* the scaled derivative.

- As we can see, the gradient ascent method does not require a second derivative
- However, it may require more iterations to converge than Newton-Raphson

Newton-Raphson approach for MLE in logistic regression

Newton-Raphson approach for MLE in logistic regression

The optimal point $\hat{\beta}$ is given by the root of the equation $\nabla_{\beta}\ell = 0$.

Newton-Raphson approach for MLE in logistic regression

The optimal point $\hat{\beta}$ is given by the root of the equation $\nabla_{\beta}\ell = 0$.

Applying Newton-Raphson, the update step is:

Newton-Raphson approach for MLE in logistic regression

The optimal point $\hat{\beta}$ is given by the root of the equation $\nabla_{\beta}\ell = 0$.

Applying Newton-Raphson, the update step is:

$$\beta_{k+1} = \beta_k - \mathbf{H}_{\beta_k}^{-1}(\ell) \nabla_{\beta_k} \ell(\beta_k) \quad (45)$$

Newton-Raphson approach for MLE in logistic regression

The optimal point $\hat{\beta}$ is given by the root of the equation $\nabla_{\beta}\ell = 0$.

Applying Newton-Raphson, the update step is:

$$\beta_{k+1} = \beta_k - \mathbf{H}_{\beta_k}^{-1}(\ell) \nabla_{\beta_k} \ell(\beta_k) \quad (45)$$

The operator \mathbf{H}^{-1} represents the inverse **Hessian** (second derivative) matrix of ℓ with respect to β :

Newton-Raphson approach for MLE in logistic regression

The optimal point $\hat{\beta}$ is given by the root of the equation $\nabla_{\beta}\ell = 0$.

Applying Newton-Raphson, the update step is:

$$\beta_{k+1} = \beta_k - \mathbf{H}_{\beta_k}^{-1}(\ell) \nabla_{\beta_k} \ell(\beta_k) \quad (45)$$

The operator \mathbf{H}^{-1} represents the inverse **Hessian** (second derivative) matrix of ℓ with respect to β :

$$\mathbf{H}_{\beta_k}(\beta_k) = \nabla_{\beta_k}^2 \ell = \begin{pmatrix} \frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} & \frac{\partial^2 \ell(\beta)}{\partial \beta_0 \beta_1} \\ \frac{\partial^2 \ell(\beta)}{\partial \beta_1 \beta_0} & \frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} \end{pmatrix} \quad (46)$$

Newton-Raphson approach for MLE in logistic regression

The optimal point $\hat{\beta}$ is given by the root of the equation $\nabla_{\beta}\ell = 0$.

Applying Newton-Raphson, the update step is:

$$\beta_{k+1} = \beta_k - \mathbf{H}_{\beta_k}^{-1}(\ell) \nabla_{\beta_k} \ell(\beta_k) \quad (45)$$

The operator \mathbf{H}^{-1} represents the inverse **Hessian** (second derivative) matrix of ℓ with respect to β :

$$\mathbf{H}_{\beta_k}(\beta_k) = \nabla_{\beta_k}^2 \ell = \begin{pmatrix} \frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} & \frac{\partial^2 \ell(\beta)}{\partial \beta_0 \beta_1} \\ \frac{\partial^2 \ell(\beta)}{\partial \beta_1 \beta_0} & \frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} \end{pmatrix} \quad (46)$$

Note that the (45) is just the matrix representation of the 1-D case:

Newton-Raphson approach for MLE in logistic regression

The optimal point $\hat{\beta}$ is given by the root of the equation $\nabla_{\beta}\ell = 0$.

Applying Newton-Raphson, the update step is:

$$\beta_{k+1} = \beta_k - \mathbf{H}_{\beta_k}^{-1}(\ell) \nabla_{\beta_k} \ell(\beta_k) \quad (45)$$

The operator \mathbf{H}^{-1} represents the inverse **Hessian** (second derivative) matrix of ℓ with respect to β :

$$\mathbf{H}_{\beta_k}(\beta_k) = \nabla_{\beta_k}^2 \ell = \begin{pmatrix} \frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} & \frac{\partial^2 \ell(\beta)}{\partial \beta_0 \beta_1} \\ \frac{\partial^2 \ell(\beta)}{\partial \beta_1 \beta_0} & \frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} \end{pmatrix} \quad (46)$$

Note that the (45) is just the matrix representation of the 1-D case:

$$\beta_{k+1} = \beta_k - \frac{\ell'(\beta_k)}{\ell''(\beta_k)} \quad (47)$$

Newton-Raphson approach for MLE (cont.)

Newton-Raphson approach for MLE (cont.)

We can work out each component of the second derivative:

Newton-Raphson approach for MLE (cont.)

We can work out each component of the second derivative:

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} = - \sum_i p(x_i)(1 - p(x_i)) \quad (48)$$

Newton-Raphson approach for MLE (cont.)

We can work out each component of the second derivative:

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} = - \sum_i p(x_i)(1 - p(x_i)) \quad (48)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0 \partial \beta_1} = - \sum_i x_i p(x_i)(1 - p(x_i)) \quad (49)$$

Newton-Raphson approach for MLE (cont.)

We can work out each component of the second derivative:

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} = - \sum_i p(x_i)(1 - p(x_i)) \quad (48)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0 \partial \beta_1} = - \sum_i x_i p(x_i)(1 - p(x_i)) \quad (49)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} = - \sum_i x_i^2 p(x_i)(1 - p(x_i)) \quad (50)$$

Newton-Raphson approach for MLE (cont.)

We can work out each component of the second derivative:

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} = - \sum_i p(x_i)(1 - p(x_i)) \quad (48)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0 \partial \beta_1} = - \sum_i x_i p(x_i)(1 - p(x_i)) \quad (49)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} = - \sum_i x_i^2 p(x_i)(1 - p(x_i)) \quad (50)$$

The complete update can then be shown as:

Newton-Raphson approach for MLE (cont.)

We can work out each component of the second derivative:

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} = - \sum_i p(x_i)(1 - p(x_i)) \quad (48)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0 \partial \beta_1} = - \sum_i x_i p(x_i)(1 - p(x_i)) \quad (49)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} = - \sum_i x_i^2 p(x_i)(1 - p(x_i)) \quad (50)$$

The complete update can then be shown as:

$$\begin{pmatrix} \beta_{0,k+1} \\ \beta_{1,k+1} \end{pmatrix} = \begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} - \left[\begin{pmatrix} \frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} & \frac{\partial^2 \ell(\beta)}{\partial \beta_0 \partial \beta_1} \\ \frac{\partial^2 \ell(\beta)}{\partial \beta_1 \partial \beta_0} & \frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} \end{pmatrix}^{-1} \begin{pmatrix} \frac{\partial \ell}{\partial \beta_0} \\ \frac{\partial \ell}{\partial \beta_1} \end{pmatrix} \right]_{\beta_k} \quad (51)$$

Newton-Raphson approach for MLE (cont.)

We can work out each component of the second derivative:

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} = - \sum_i p(x_i)(1 - p(x_i)) \quad (48)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0 \beta_1} = - \sum_i x_i p(x_i)(1 - p(x_i)) \quad (49)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} = - \sum_i x_i^2 p(x_i)(1 - p(x_i)) \quad (50)$$

The complete update can then be shown as:

$$\begin{pmatrix} \beta_{0,k+1} \\ \beta_{1,k+1} \end{pmatrix} = \begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} - \left[\begin{pmatrix} \frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} & \frac{\partial^2 \ell(\beta)}{\partial \beta_0 \partial \beta_1} \\ \frac{\partial^2 \ell(\beta)}{\partial \beta_1 \partial \beta_0} & \frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} \end{pmatrix}^{-1} \begin{pmatrix} \frac{\partial \ell}{\partial \beta_0} \\ \frac{\partial \ell}{\partial \beta_1} \end{pmatrix} \right]_{\beta_k} \quad (51)$$

Alternatively:

Newton-Raphson approach for MLE (cont.)

We can work out each component of the second derivative:

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} = - \sum_i p(x_i)(1 - p(x_i)) \quad (48)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_0 \beta_1} = - \sum_i x_i p(x_i)(1 - p(x_i)) \quad (49)$$

$$\frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} = - \sum_i x_i^2 p(x_i)(1 - p(x_i)) \quad (50)$$

The complete update can then be shown as:

$$\begin{pmatrix} \beta_{0,k+1} \\ \beta_{1,k+1} \end{pmatrix} = \begin{pmatrix} \beta_{0,k} \\ \beta_{1,k} \end{pmatrix} - \left[\begin{pmatrix} \frac{\partial^2 \ell(\beta)}{\partial \beta_0^2} & \frac{\partial^2 \ell(\beta)}{\partial \beta_0 \partial \beta_1} \\ \frac{\partial^2 \ell(\beta)}{\partial \beta_1 \partial \beta_0} & \frac{\partial^2 \ell(\beta)}{\partial \beta_1^2} \end{pmatrix}^{-1} \begin{pmatrix} \frac{\partial \ell}{\partial \beta_0} \\ \frac{\partial \ell}{\partial \beta_1} \end{pmatrix} \right]_{\beta_k} \quad (51)$$

Alternatively:

$$\beta_{k+1} = \beta_k - \left[\left(\frac{\partial^2 \ell(\beta)}{\partial \beta \partial \beta^T} \right)^{-1} \frac{\partial \ell(\beta)}{\partial \beta} \right]_{\beta_k} \quad (52)$$

Constrained optimization

Constrained optimization

Deals with problems where we seek to minimize an objective:

Constrained optimization

Deals with problems where we seek to minimize an objective:

$$\min_{\theta \in \mathcal{C}} \mathcal{L}(\theta) \quad (53)$$

Constrained optimization

Deals with problems where we seek to minimize an objective:

$$\min_{\theta \in \mathcal{C}} \mathcal{L}(\theta) \quad (53)$$

subject to:

Constrained optimization

Deals with problems where we seek to minimize an objective:

$$\min_{\theta \in \mathcal{C}} \mathcal{L}(\theta) \quad (53)$$

subject to:

$$h_i(\theta) = 0, \quad i \in \mathcal{E} \quad (54)$$

$$g_j(\theta) \leq 0, \quad j \in \mathcal{I} \quad (55)$$

where:

- \mathcal{C} : constraint/feasible set
- \mathcal{E} : set of equality constraints
- \mathcal{I} : set of inequality constraints

Constrained optimization

Deals with problems where we seek to minimize an objective:

$$\min_{\theta \in \mathcal{C}} \mathcal{L}(\theta) \quad (53)$$

subject to:

$$h_i(\theta) = 0, \quad i \in \mathcal{E} \quad (54)$$

$$g_j(\theta) \leq 0, \quad j \in \mathcal{I} \quad (55)$$

Constrained optimization

Deals with problems where we seek to minimize an objective:

$$\min_{\theta \in \mathcal{C}} \mathcal{L}(\theta) \quad (53)$$

subject to:

$$h_i(\theta) = 0, \quad i \in \mathcal{E} \quad (54)$$

$$g_j(\theta) \leq 0, \quad j \in \mathcal{I} \quad (55)$$

where:

- \mathcal{C} : constraint/feasible set

Constrained optimization

Deals with problems where we seek to minimize an objective:

$$\min_{\theta \in \mathcal{C}} \mathcal{L}(\theta) \quad (53)$$

subject to:

$$h_i(\theta) = 0, \quad i \in \mathcal{E} \quad (54)$$

$$g_j(\theta) \leq 0, \quad j \in \mathcal{I} \quad (55)$$

where:

- \mathcal{C} : constraint/feasible set
- \mathcal{E} : set of equality constraints

Constrained optimization

Deals with problems where we seek to minimize an objective:

$$\min_{\theta \in \mathcal{C}} \mathcal{L}(\theta) \quad (53)$$

subject to:

$$h_i(\theta) = 0, \quad i \in \mathcal{E} \quad (54)$$

$$g_j(\theta) \leq 0, \quad j \in \mathcal{I} \quad (55)$$

where:

- \mathcal{C} : constraint/feasible set
- \mathcal{E} : set of equality constraints
- \mathcal{I} : set of inequality constraints

Lagrange multipliers

Lagrange multipliers

Given an optimization problem with m equality constraints $\mathbf{h}(\boldsymbol{\theta})$, we can write the Lagrangian as:

Lagrange multipliers

Given an optimization problem with m equality constraints $\mathbf{h}(\boldsymbol{\theta})$, we can write the Lagrangian as:

$$\mathcal{L}(\boldsymbol{\theta}, \lambda) := \mathcal{L}(\boldsymbol{\theta}) + \sum_{j=1}^m \lambda_j h_j(\boldsymbol{\theta}) \quad (56)$$

Lagrange multipliers

Given an optimization problem with m equality constraints $\mathbf{h}(\boldsymbol{\theta})$, we can write the Lagrangian as:

$$\mathcal{L}(\boldsymbol{\theta}, \lambda) := \mathcal{L}(\boldsymbol{\theta}) + \sum_{j=1}^m \lambda_j h_j(\boldsymbol{\theta}) \quad (56)$$

where λ_j are the Lagrange multipliers.

Lagrange multipliers

Given an optimization problem with m equality constraints $\mathbf{h}(\boldsymbol{\theta})$, we can write the Lagrangian as:

$$\mathcal{L}(\boldsymbol{\theta}, \lambda) := \mathcal{L}(\boldsymbol{\theta}) + \sum_{j=1}^m \lambda_j h_j(\boldsymbol{\theta}) \quad (56)$$

where λ_j are the Lagrange multipliers.

Thus, to find $*$, we solve:

Lagrange multipliers

Given an optimization problem with m equality constraints $\mathbf{h}(\boldsymbol{\theta})$, we can write the Lagrangian as:

$$\mathcal{L}(\boldsymbol{\theta}, \lambda) := \mathcal{L}(\boldsymbol{\theta}) + \sum_{j=1}^m \lambda_j h_j(\boldsymbol{\theta}) \quad (56)$$

where λ_j are the Lagrange multipliers.

Thus, to find $*$, we solve:

$$\nabla_{\boldsymbol{\theta}, \lambda} \mathcal{L}(\boldsymbol{\theta}, \lambda) = \mathbf{0} \quad (57)$$

Lagrange multipliers

Given an optimization problem with m equality constraints $\mathbf{h}(\boldsymbol{\theta})$, we can write the Lagrangian as:

$$\mathcal{L}(\boldsymbol{\theta}, \lambda) := \mathcal{L}(\boldsymbol{\theta}) + \sum_{j=1}^m \lambda_j h_j(\boldsymbol{\theta}) \quad (56)$$

where λ_j are the Lagrange multipliers.

Thus, to find $*$, we solve:

$$\nabla_{\boldsymbol{\theta}, \lambda} \mathcal{L}(\boldsymbol{\theta}, \lambda) = \mathbf{0} \quad (57)$$

- If $\boldsymbol{\theta} \in \mathbb{R}^D$, there are $D + m$ equations with an equal number of unknowns

Karush-Kuhn-Tucker (KKT) conditions

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) +$$

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) + \sum_i \mu_i g_i(\boldsymbol{\theta}) + \sum_j \lambda_j h_j(\boldsymbol{\theta}) \quad (58)$$

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) + \sum_i \mu_i g_i(\boldsymbol{\theta}) + \sum_j \lambda_j h_j(\boldsymbol{\theta}) \quad (58)$$

The optimization problem then becomes:

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) + \sum_i \mu_i g_i(\boldsymbol{\theta}) + \sum_j \lambda_j h_j(\boldsymbol{\theta}) \quad (58)$$

The optimization problem then becomes:

$$\min_{\boldsymbol{\theta}} \max_{\boldsymbol{\mu} \geq \mathbf{0}, \boldsymbol{\lambda}} L(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) \quad (59)$$

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) + \sum_i \mu_i g_i(\boldsymbol{\theta}) + \sum_j \lambda_j h_j(\boldsymbol{\theta}) \quad (58)$$

The optimization problem then becomes:

$$\min_{\boldsymbol{\theta}} \max_{\boldsymbol{\mu} \geq \mathbf{0}, \boldsymbol{\lambda}} L(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) \quad (59)$$

When \mathcal{L} and g are convex, then the KKT conditions are necessary and sufficient for global optimality:

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) + \sum_i \mu_i g_i(\boldsymbol{\theta}) + \sum_j \lambda_j h_j(\boldsymbol{\theta}) \quad (58)$$

The optimization problem then becomes:

$$\min_{\boldsymbol{\theta}} \max_{\boldsymbol{\mu} \geq \mathbf{0}, \boldsymbol{\lambda}} L(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) \quad (59)$$

When \mathcal{L} and g are convex, then the KKT conditions are necessary and sufficient for global optimality:

- Feasibility: constraints satisfied

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) + \sum_i \mu_i g_i(\boldsymbol{\theta}) + \sum_j \lambda_j h_j(\boldsymbol{\theta}) \quad (58)$$

The optimization problem then becomes:

$$\min_{\boldsymbol{\theta}} \max_{\boldsymbol{\mu} \geq \mathbf{0}, \boldsymbol{\lambda}} L(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) \quad (59)$$

When \mathcal{L} and g are convex, then the KKT conditions are necessary and sufficient for global optimality:

- Feasibility: constraints satisfied
- Stationarity of solution: $\nabla_{\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}} L = \mathbf{0}$

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) + \sum_i \mu_i g_i(\boldsymbol{\theta}) + \sum_j \lambda_j h_j(\boldsymbol{\theta}) \quad (58)$$

The optimization problem then becomes:

$$\min_{\boldsymbol{\theta}} \max_{\boldsymbol{\mu} \geq \mathbf{0}, \boldsymbol{\lambda}} L(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) \quad (59)$$

When \mathcal{L} and g are convex, then the KKT conditions are necessary and sufficient for global optimality:

- Feasibility: constraints satisfied
- Stationarity of solution: $\nabla_{\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}} L = \mathbf{0}$
- Dual feasibility: $\boldsymbol{\mu} \geq \mathbf{0}$

Karush-Kuhn-Tucker (KKT) conditions

Given a problem with equality constraints $\mathbf{h}(\boldsymbol{\theta}) = \mathbf{0}$ and inequality constraints $\mathbf{g}(\boldsymbol{\theta}) \leq \mathbf{0}$, the generalized Lagrangian is given by:

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) = \mathcal{L}(\boldsymbol{\theta}) + \sum_i \mu_i g_i(\boldsymbol{\theta}) + \sum_j \lambda_j h_j(\boldsymbol{\theta}) \quad (58)$$

The optimization problem then becomes:

$$\min_{\boldsymbol{\theta}} \max_{\boldsymbol{\mu} \geq \mathbf{0}, \boldsymbol{\lambda}} L(\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}) \quad (59)$$

When \mathcal{L} and g are convex, then the KKT conditions are necessary and sufficient for global optimality:

- Feasibility: constraints satisfied
- Stationarity of solution: $\nabla_{\boldsymbol{\theta}, \boldsymbol{\mu}, \boldsymbol{\lambda}} L = \mathbf{0}$
- Dual feasibility: $\boldsymbol{\mu} \geq \mathbf{0}$
- Complementary slackness: $\boldsymbol{\mu} \odot \mathbf{g} = \mathbf{0}$

Further topics

Some to be covered in Advanced Probabilistic ML:

Further topics

Some to be covered in Advanced Probabilistic ML:

- Linear programming (simplex algorithm)
- Quadratic programming
- Proximal gradient method
- Bound optimization (majorize-minimize algorithms)
- Expectation maximization (EM); for MLE/MAP estimation

Reading assignments

- **PMLI 8**
- **PMLCE 5**