

CEE 697M: Probabilistic Machine Learning

M4 Nonparametric Methods:

L4a: Exemplar-based methods

Jimi Oke

UMass**Amherst**

College of Engineering

Wed, Apr 19, 2023

Outline

- 1 Introduction
- 2 KNN
- 3 Metric learning
- 4 Density kernels
- 5 Kernel smoothing
- 6 Local regression
- 7 Outlook

Nonparametric modeling

Parametric models seek to estimate $p(\mathbf{y}|\boldsymbol{\theta})$ (unconditional case) or $p(\mathbf{y}|\mathbf{x}, \boldsymbol{\theta})$ (conditional case).

- $\boldsymbol{\theta}$ is a fixed-dimensional vector of **parameters**
- Estimation is performed using a dataset $\mathcal{D} = \{(\mathbf{x}_n, \mathbf{y}_n) : n = 1 : N\}$
- There is an assumed functional form: $\mathbf{y} \sim f_{\boldsymbol{\theta}}(\mathbf{x})$

Nonparametric models are defined based on similarity between a test input \mathbf{x} at each training input \mathbf{x}_n : $d(\mathbf{x}, \mathbf{x}_n)$

- No assumption of functional form on model parameters
- Effective number of parameters can grow with size of dataset $|\mathcal{D}|$
- Known as **exemplar-based models** (as training samples are used to make each future prediction)
- Other names: instance-based learning, memory-based learning

Exemplar-based models

We will consider the following exemplar approaches:

- K-nearest neighbors (KNN)
- Kernel density estimation
- Kernel [local] regression

K nearest neighbor classifier

Basic idea: classify new/test input \mathbf{x} by assigning to most probable (majority) label in the neighborhood of \mathbf{x} (closest examples) from the training set.

Thus, we estimate:

$$p(y = c | \mathbf{x}, \mathcal{D}) = \frac{1}{K} \sum_{n \in N_K(\mathbf{x}, \mathcal{D})} \mathbb{I}(y_n = c) \quad (1)$$

- c class label
- K : number of training samples in neighborhood
- $N_K(\mathbf{x}, \mathcal{D})$: neighborhood of \mathbf{x} (size K) based on dataset \mathcal{D}

Illustration of KNN




Figure: Illustration of the KNN approach on a training set of 12 observations and the resulting decision boundary. (ESL Fig 2.14)

Bias-variance trade-off in KNN

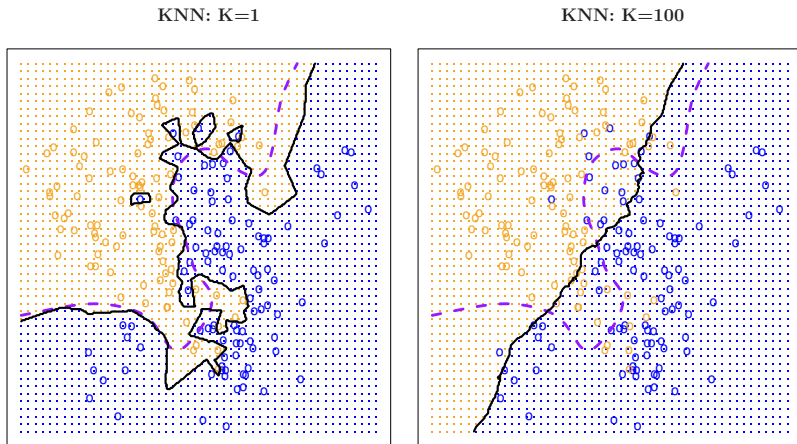


Figure: Comparing decision boundaries $K = 1$ and $K = 100$ for a dataset of 100 observations. Which model has lower bias? Which one gives a higher variance? The Bayes decision boundary is the purple dashed line (ESL Fig 2.16)

Approximating Bayes decision boundary with KNN




Figure: KNN decision boundary with $K = 10$ on the same training data set. (ESL Fig 2.15)

Training and test error rates for KNN





Figure: KNN training error rate (blue, 200 observations) and test error rate (orange, 5000 observations). Flexibility increases as K decreases. Which K should you choose? (ESL Fig 2.17)

KNN considerations

- To find the points in the neighborhood N_K , we need to determine the K -closest points to input \mathbf{x} . This is done via a specified distance metric: $d(\mathbf{x}, \mathbf{x}') \in \mathbb{R}^+$ (e.g. Euclidean, Mahalanobis)
- $K = 1$ induces a Voronoi tessellation: partitioning of input space such that all points $\mathbf{x} \in V(\mathbf{x}_n)$ are closer to \mathbf{x}_n than to any other point (From a modeling perspective, this is overfitting)




Source: <https://package.elm-lang.org/packages/ianmackenzie/elm-geometry/latest/VoronoiDiagram2d>

- Suffers under high dimensionality
- Memory intensive

KNN extension: open set recognition

KNN is readily applicable to open set recognition (set of classes \mathcal{C} not fixed).

- Novelty/out-of-distribution/anomaly detection



Source: <https://www.intechopen.com/chapters/74393>

- Incremental/online/life-long/continual learning: any potentially new label is added to new class C_{t+1} ; dataset augmented
- Few-shot classification (for person re-identification or face verification)

Distance metrics

The [semantic] distance between points \mathbf{x} and \mathbf{x}' is specified by a distance metric $d(\mathbf{x}, \mathbf{x}') \in \mathbb{R}^+$.

- Alternately, the similarity $s(\mathbf{x}, \mathbf{x}')$ can be computed
- Distance/similarity required for KNN, unsupervised learning (e.g. clustering) among other tasks
- Common metrics:
 - Euclidean distance:

$$d_E(\mathbf{x}, \mathbf{x}') = \sqrt{(\mathbf{x} - \mathbf{x}')^\top (\mathbf{x} - \mathbf{x}')} \quad (2)$$

- Mahalanobis distance:

$$d_M(\mathbf{x}, \mathbf{x}') = \sqrt{(\mathbf{x} - \mathbf{x}')^\top \mathbf{M} (\mathbf{x} - \mathbf{x}')} \quad (3)$$

where \mathbf{M} is the Mahalanobis distance matrix

Distance metrics

The process of finding the optimal \mathbf{M} is called **metric learning**

- When D is large, we typically learn an embedding (mapping): $\mathbf{e} = f(\mathbf{x})$ and then compute $d_{\mathbf{M}}(\mathbf{e}, \mathbf{e}')$ instead.
- When f is a deep neural network, this process is termed **deep metric learning**

Methods for estimating M

- Large margin nearest neighbors (LMNN)
- Neighborhood component analysis (NCA)
- Latent coincidence analysis (LCA)
- Minimization of classification and ranking losses (with mining techniques and proxy methods):
 - Pairwise/contrastive loss
 - Triplet loss
 - N-pairs loss

Density kernel

A density kernel $\mathcal{K}(x)$ is a weighting function that specifies a mapping or transformation of an input x : $\mathcal{K} : \mathbb{R} \rightarrow \mathbb{R}_+$.

Density kernels have two important properties:

- **Normalization:**

$$\int x\mathcal{K}(x)dx = 0 \quad (4)$$

- **Symmetry:**

$$\mathcal{K}(-x) = \mathcal{K}x \quad (5)$$

Kernels have several uses, e.g.

- density function estimation
- local regression (our focus)
- smoothing time series

Kernel functions




Figure: Popular kernel functions

- Triangular: $\mathcal{K}(x) = (1 - |x|)$
- Epanechnikov: $\mathcal{K}(x) = \frac{3}{4}(1 - x^2)$
- Triweight: $\mathcal{K}(x) = \frac{35}{32}(1 - x^2)^3$
- Tricube: $\mathcal{K}(x) = \frac{70}{81}(1 - |x|^3)^3$
- Gaussian: $\mathcal{K}(x) = \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}x^2}$

Radial basis function

The radial basis function (RBF) kernel is a generalization of a density kernel to an input vector \mathbf{x} :

$$\mathcal{K}_h(\mathbf{x}) \propto \mathcal{K}_h(\|\mathbf{x}\|) \quad (6)$$

where h is the bandwidth parameter:

The RBF Gaussian kernel is thus given by:

$$\mathcal{K}_h(\mathbf{x}) = \frac{1}{h^D (2\pi)^{D/2}} \prod_{d=1}^D \exp \left[-\frac{1}{2h^2} x_d^2 \right] \quad (7)$$

Bandwidth parameter

Specifies the width of the kernel:

$$\mathcal{K}_h := \frac{1}{h} \mathcal{K} \left(\frac{\mathbf{x}}{h} \right) \quad (8)$$

K-nearest neighbor smoother

In the simple case of the KNN kernel, we use the neighborhood average:

$$\hat{f}(x_0) = \frac{1}{k} \sum_{i \in N_K(x_0)} y_i \tag{9}$$

where $N_K(x_0)$ is the K -nearest neighborhood of x_0 .

- All points in the neighborhood are equally weighted
- The resulting $\hat{f}(x)$ is not smooth
- To achieve smoothness, we weight observations by distance to the target point




Figure: KNN equally-weighted kernel

Nadaraya-Watson smoother

Nadaraya-Watson kernel-weighted average implements distance-based weighting:

$$\mathbb{E}[y|\mathbf{x}, \mathcal{D}] = \hat{f}(x_0) = \frac{\sum_{i=1}^n K_{\lambda}(x_0, x_i) y_i}{\sum_{i=1}^n K_{\lambda}(x_0, x_i)} \quad (10)$$

where K_{λ} can be any kernel function.

If we use the popular **Epanechnikov** (quadratic) kernel, then:

$$K_{\lambda}(x_0, x_i) = D\left(\frac{|x_i - x_0|}{\lambda}\right) \quad (11)$$

where

$$D(t) = \begin{cases} \frac{3}{4} (1 - t^2) & |t| \leq 1 \\ 0 & \text{otherwise} \end{cases} \quad (12)$$

In D , the half-width (or *bandwidth*) of the neighborhood is given by λ .

Nadaraya-Watson smoother (Epanechnikov kernel)




Figure: Nadaraya-Watson estimate of $\hat{f}(x)$ using the Epanechnikov kernel. Half-width is fixed at $\lambda = 0.2$

The half-width can be generalized as any function h_λ of the target point x_0 :

$$K_\lambda(x_0, x_i) = D \left(\frac{|x_i - x_0|}{h_\lambda(x_0)} \right) \quad (13)$$

Kernels as a localization device

Rather than estimate a regression function $f(X)$ over the entire \mathbb{R}^D , we can estimate the response at each training point using a weighted average:

- only observations close a target point x_0 are used in estimating f at that point
- the function defining how the points in the neighborhood of x_0 are weighted is called a **kernel**: $K_\lambda(x_0, x_i)$ where λ controls the width of the neighborhood
- this is a localized/memory-based approach
- the resulting $\hat{f}(X)$ is smooth in \mathbb{R}^D

Using kernel functions, we can estimate $\hat{f}(X)$ in two ways:

- ① **Nonparametric**: define an averaging function (kernel) to estimate y_0 for each point x_0
- ② **Parametric**: estimate a linear model for each point x_0

Local linear regression

Using a nonparametric average function produces biased estimates at the boundaries.




Figure: Locally weighted average (Nadaraya-Watson) versus local linear regression.

To correct this, we can estimate a linear model at each point x_0 by solving a weighted least squares:

$$\min_{\alpha(x_0), \beta(x_0)} \sum_{i=1}^n K_{\lambda}(x_0, x_i) [y_i - \alpha(x_0) - \beta(x_0)x_i]^2 \quad (14)$$

Local linear regression estimates

Let $b(x)^T = (1, x)$ and:

$$B = \begin{pmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{pmatrix} \quad W(x_0) = \begin{pmatrix} K_\lambda(x_0, x_1) & 0 & \cdots & 0 \\ 0 & K_\lambda(x_0, x_2) & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & K_\lambda(x_0, x_n) \end{pmatrix} \quad (15)$$

Then the solution to the locally weighted regression problem is:

$$\hat{f}(x_0) = b(x_0)^T (B^T W(x_0) B)^{-1} B^T W(x_0) y \quad (16)$$

$$\hat{f}(x_0) = \sum_{i=1}^n \ell_i(x_0) y_i \quad (17)$$

The weights $\ell_i(x_0)$ are called the **equivalent kernel**

Effect of equivalent kernel

The equivalent kernel (local regression) corrects the bias from local average kernel methods to the first order.

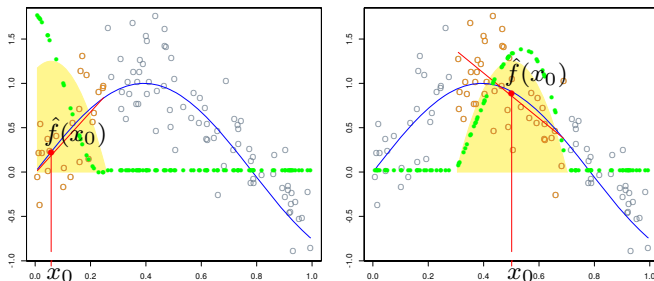


Figure: Green points show the weights $\ell_i(x_0)$ (rescaled for display purposes) along with those for the Nadaraya-Watson (N-W) local average (yellow shaded region; also rescaled). The correction effect of local regression can be observed.

Local polynomial regression

- Higher-order terms in $\hat{f}(x)$ are required to reduce bias in curved regions
- Local polynomial regression can correct this at the cost of higher variance
- Given by:

$$\hat{f}(x_0) = \hat{\alpha}(x_0) + \sum_{j=1}^d \hat{\beta}_j(x_0) x_0^j \quad (18)$$

where $\hat{f}(x_0)$ is the solution to:

$$\min_{\alpha(x_0), \beta(x_0), j=1, \dots, d} \sum_{i=1}^n K_{\lambda}(x_0, x_i) \left[y_i - \alpha(x_0) - \sum_{j=1}^d \beta_j(x_0) x_i^j \right]^2 \quad (19)$$

Bias correction of local quadratic regression

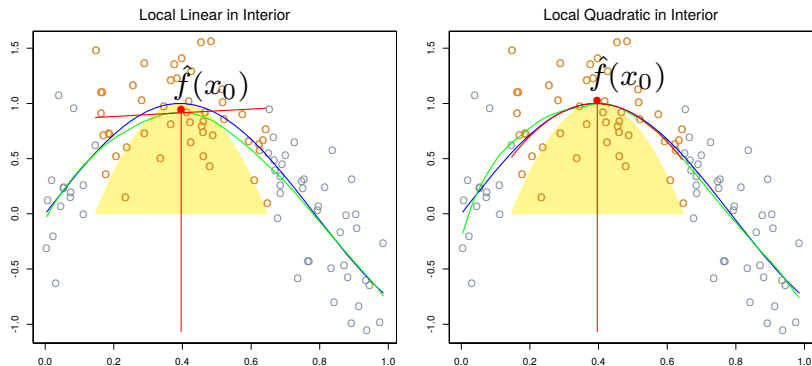


Figure: The local linear regression estimator is biased in curved regions. A higher-order local fit (in this case, quadratic) can eliminate this.

Local regression algorithms

- LOESS: locally estimated scatterplot smoothing
 - Identical approach earlier developed in 1964 by [Savitzky and Golay](#) for smoothing noisy data (known as the Savitzky-Golay filter)
 - “Rediscovered” by [William Cleveland in 1979](#)
- LOWESS: locally weighted scatterplot smoothing
 - Extension of LOESS by [Cleveland and Susan Devlin \(1988\)](#)
- LOESS can be considered a generalization of LOWESS, as it fits multivariate data, while LOWESS is for univariate cases

Summary

- Kernel smoothing methods can be used for flexible functional fitting
- Local regression generates a linear/polynomial fit at each target point using a kernel weighted loss function

Reading:

- **PMLI 16**
- **ESL 6.1** One-dimensional Kernel Smoothers (pp. 191–199)
- **ISLR 7.6:** Local Regression (pp. 280–282)