

CEE 616: Probabilistic Machine Learning  
M4 Nonparametric Methods:  
L4D: Trees and Ensemble Methods

**Jimi Oke**

UMass**Amherst**

---

College of Engineering

Tue, Nov 18, 2025

# Outline

① Introduction

② Regression trees

③ Classification trees

④ Bagging

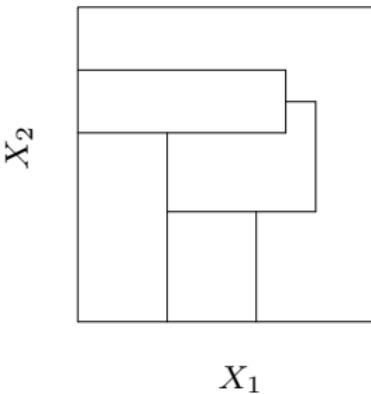
⑤ Random Forests

⑥ Boosting

⑦ Summary

# Decision trees

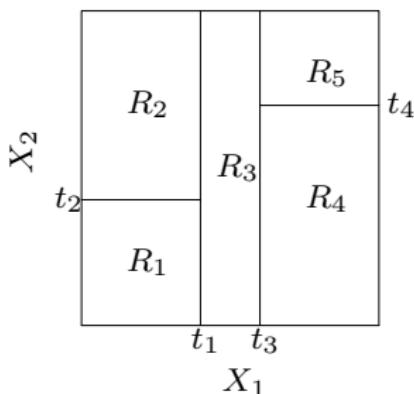
- Consider a regression problem with continuous response  $y$  and inputs  $X_1$  and  $X_2$ .
- Now, imagine we want to model response  $y$  as a constant across different regions of the input space.
- One option can be:



- Can you describe the region using inequalities?
- How would you estimate response  $y$  in each region?

# Recursive binary partitioning

In order to ensure simplicity and interpretability, we consider splitting a dataset using the **recursive binary partitioning** (RBP) algorithm

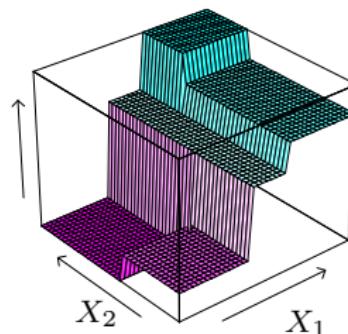
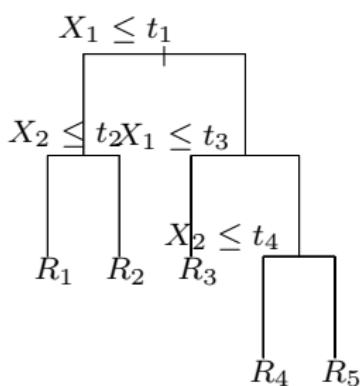


Steps:

- ① Split the input space at  $X_1 = t_1$ ; the response is modeled as the mean of  $y$  in each region
- ② Split the region  $X_1 \leq t_1$  at  $X_2 = t_2$  and the region  $X_1 > t_1$  at  $X_1 = t_3$
- ③ Split the region  $X_1 > t_3$  at  $X_2 = t_4$

Thus we obtain 5 regions. In each case, the *decision* is to choose the **variable** and the **split-point** that provide the best fit.

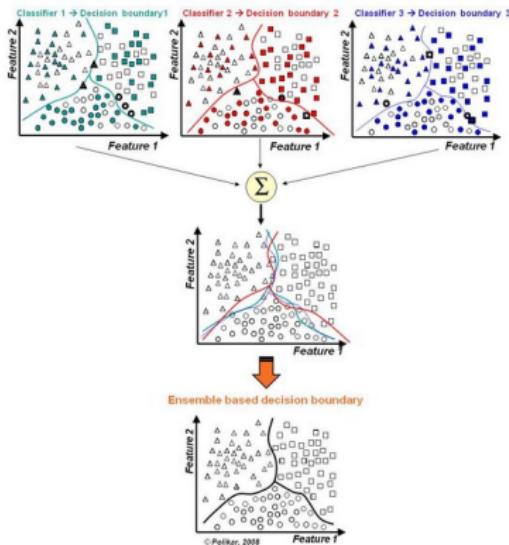
# Decision trees (cont.)



- The partitioning steps can be displayed as a **decision tree**
- The split points are the **internal nodes**
- Lines connecting the nodes are called *branches*
- The final regions are **terminal nodes** or **leaves** ( $R_1$ ,  $R_2$ ,  $R_3$ ,  $R_4$  and  $R_5$  in the figure)

# Ensemble learning

Ensemble methods employ various techniques to combine predictions from multiple models. This approach can be termed “learning by committee”



We will consider three approaches with a focus on decision trees:

- ① **Bagging:** average/majority vote over bootstrap sample
- ② **Random Forests:** average/majority vote over collection of de-correlated trees
- ③ **Boosting:** weighted aggregation of weak learners

Source: [http://www.scholarpedia.org/article/File:Combining\\_classifiers2.jpg](http://www.scholarpedia.org/article/File:Combining_classifiers2.jpg)

# Regression tree

Given  $D$  features and  $N$  observations, how do we predict the response for an observation  $x_0$  using a regression tree?

The goal is to partition the input space into  $M$  regions using the **recursive binary partitioning** (RBP) approach.

- This is a greedy, top-down approach
- We find the RSS-minimizing split (variable-cutpoint pair) for each subsequent region
- Tree-growing is terminated using a reasonable stopping condition

# Recursive binary partitioning decision

For a given  $j$ th variable and split point  $s$ , we have a pair of half-planes:

$$R_1(j, s) = \{\mathbf{x} : x_j < t\} \quad (1)$$

$$R_2(j, s) = \{\mathbf{x} : x_j \geq t\} \quad (2)$$

Find splitting variable  $j^*$  and point  $s^*$  that solve:

$$j^*, s^* = \arg \min_{j,t} \left[ \sum_{x_i \in R_1(j,t)} (y_i - \hat{y}_{R_1})^2 + \sum_{x_i \in R_2(j,t)} (y_i - \hat{y}_{R_2})^2 \right] \quad (3)$$

where

$$\hat{y}_{R_1} = \text{ave}(y_i : \mathbf{x}_i \in R_1(j, t)) \quad (4)$$

$$\hat{y}_{R_2} = \text{ave}(y_i : \mathbf{x}_i \in R_2(j, t)) \quad (5)$$

# Overfitting and pruning

- Earlier, we mentioned tree partitioning ends when the appropriate terminating condition is met
  - Traditional thresholding can be myopic, as a potentially better solutions could be discarded.
- Also, we want to avoid excessively large trees, as these are prone to overfitting
- How then can we guarantee the best result?

## Strategy

- Grow a large tree  $T_0$  to a minimum specified node size
- Use *cost-complexity pruning* to prune  $T_0$  to find an acceptable subtree  $T \subset T_0$

# Cost-complexity pruning

Let  $m$  be the index of terminal nodes (recall that each node represents a region). Then let:

$$N_m = \sum_{i=1}^N \mathbb{I}(\mathbf{x}_i \in R_m) \quad (6)$$

$$\hat{y}_m = \frac{1}{N_m} \sum_{\mathbf{x}_i \in R_m} y_i \quad (7)$$

$$Q_m(T) = \frac{1}{N_m} \sum_{\mathbf{x}_i \in R_m} (y_i - \hat{y}_m)^2 \quad (\text{a measure of impurity}) \quad (8)$$

The **cost-complexity criterion** is then given by:

$$C_\alpha(T) = \sum_{m=1}^{|T|} N_m Q_m(T) + \alpha |T| \quad (9)$$

where  $|T|$  is the number of terminal nodes in  $T$  and  $\alpha \geq 0$  is a tuning parameter.

# Cost-complexity pruning (cont.)

Using the cost-complexity criterion, we find subtree  $T_\alpha \subseteq T_0$  that minimizes  $C_\alpha(T)$ .

- $\alpha|T|$  is a complexity penalty term
- If  $\alpha = 0$ , then  $T_\alpha = T_0$
- Larger  $\alpha$  values result in smaller trees  $T_\alpha$
- We find  $T_\alpha$  using a procedure known as *weakest link pruning*<sup>1</sup>

Your final question should be, how do we find the optimal  $\hat{\alpha}$ ?

## Answer

Cross-validation ( $k$ -fold)

<sup>1</sup>Consult ESL p. 308 for further details on this algorithm.

# Classification trees

- In regression trees, the predicted response for an observation in a certain region (terminal node) is given by the *mean response* of observations at that node
- In classification trees, the *most commonly occurring class* is the predicted response
- To make the binary splits in classification trees, we use a suitable impurity measure in place of the RSS

The proportion of class  $c$  observations in node  $m$  is given by:

$$\hat{p}_{mc} = \frac{1}{N_m} \sum_{x_i \in R_m} \mathbb{I}(y_i = c) \quad (10)$$

We then classify observations in node  $m$  to the class with maximum representation:

$$c(m) = \arg \max_k \hat{p}_{mc} \quad (11)$$

# Node impurity measures

The recursive binary partitioning splits the observations into regions  $R_m$  to minimize **node impurity** at  $m$ , measured by:

- **Misclassification error**

$$E = \frac{1}{N_m} \sum_{i \in R_m} \mathbb{I}(y_i \neq c(m)) = 1 - \hat{p}_{mc(m)} \quad (12)$$

where  $ck(m) = \arg \max_k \hat{p}_{mc}$

- **Gini index**

$$G = \sum_{c \neq c'} \hat{p}_{mc} \hat{p}_{mc'} = \sum_{c=1}^C \hat{p}_{mc} (1 - \hat{p}_{mc}) \quad (13)$$

- **Cross-entropy (deviance)**

$$D = - \sum_{c=1}^C \hat{p}_{mc} \log \hat{p}_{mc} \quad (14)$$

# Node impurity measures (cont.)

- The Gini index and the cross-entropy are more sensitive to changes in node probabilities
- They are also continuously differentiable can be readily optimized

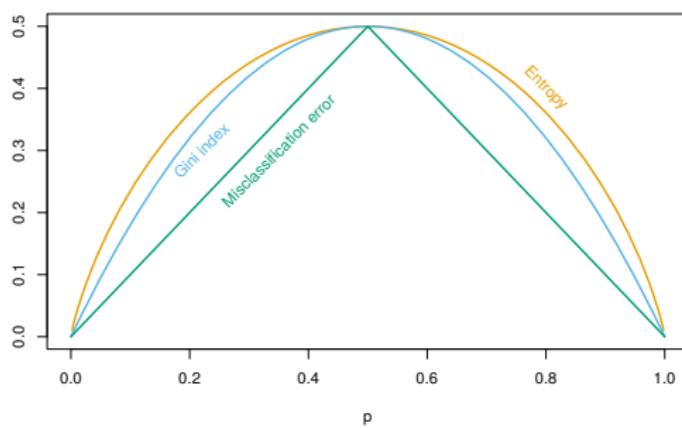


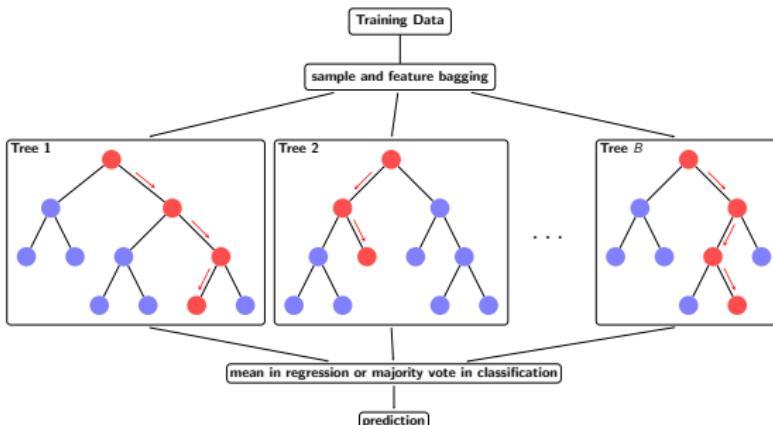
Figure: Various impurity measures plotted against  $p$  in the two-class case

Typically, the misclassification error is used in pruning, while the others are used in growing the tree

# BAGGING: Bootstrap AGGREGATING

Consider a training data set  $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$ . Let a fitted decision tree be given by  $\hat{f}(\mathbf{x})$ .

- Predictions have low bias but high variance (trees are unstable)
- Obtaining predictions from trees fitted to  $B$  bootstrap samples reduces variance (**bagging**)



# Bagging algorithm

- ① Sample with replacement from  $\mathcal{D}$  to generate  $B$  bootstrap samples,  
 $\{\mathcal{D}^{*b} : b = 1, 2, \dots, B\}$
- ② Fit a decision tree to each  $\mathcal{D}^{*b}$  to obtain prediction  $\hat{f}^{*b}(x)$
- ③ Obtain the **bagging estimate**:
  - a Regression: Average the predictions

$$\hat{f}_{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x) \quad (15)$$

- b** Classification: majority vote

$$\hat{f}_{\text{bag}}(x) = \arg \max_c \sum_{b=1}^B \mathbb{I}(\hat{f}^{*b}(x) = c) \quad (16)$$

# Example 1: Bagging trees on simulated dataset

Given a dataset of  $n = 30$  observations with  $p = 5$  predictors

- Pairwise correlation for each predictor is 0.95
- 200 bootstrap samples are created and unpruned decision trees are fitted
- High correlations leads to high variance in individual decision tree estimates

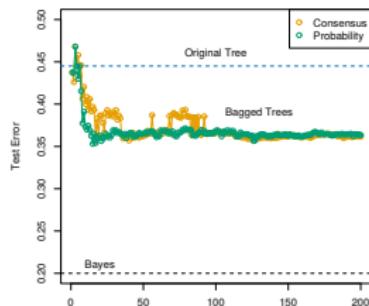
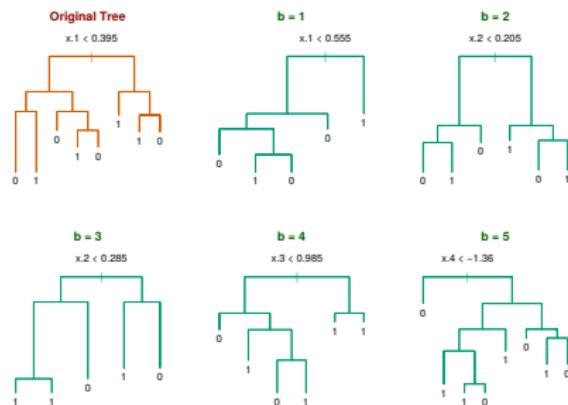
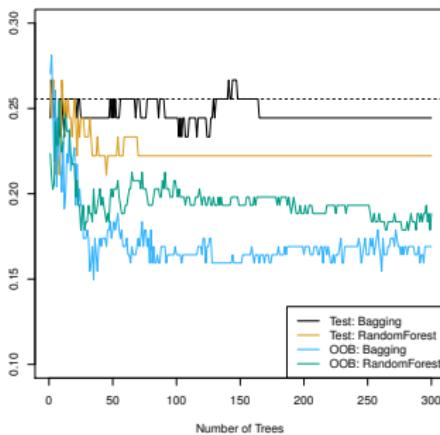


Figure: (L): Fitted trees on 6 bootstrap samples. (R): Test errors.

# Out-of-bag error

Bootstrap samples only contain 63% of the unique observations on average. Why?

- Observations left out of bootstrap sample are **out-of-bag** (OOB) instances
- Errors can be computed on each OOB instance per bootstrap sample
- These are then averaged to obtain a validation error estimate for use in hyperparameter selection



**Figure:** Comparing test (validation set approach) and OOB (similar to cross-validation) errors on bagging and random forest estimates on the Heart dataset (survival of heart transplant patients).

# Feature importance

Bagging provides stability (reduced variance) at the expense of simplicity (interpretability).

- To aid inference, we compute **feature importances** of bagged tree predictors:

$$fi_d(T) = \sum_{m=1}^{M-1} G_m \mathbb{I}(t_m = d) \quad (17)$$

where

- $fi_d(T)$  is the importance of feature  $d$  in tree  $T$
- $G_m$  is the gain in accuracy (reduction in loss) at node  $m$
- $j_m = d$  if node  $m$  uses feature  $d$  as its splitting variable  $t$
- These are then averaged over all  $B$  trees (estimators)
- Regression trees: total amount of RSS decreases due to splits
- Classification trees: Gini index used

# Random forests

Invented by [Leo Breiman in 2001](#), **random forests** build on bagging by fitting *decorrelated* trees on the bootstrap samples.

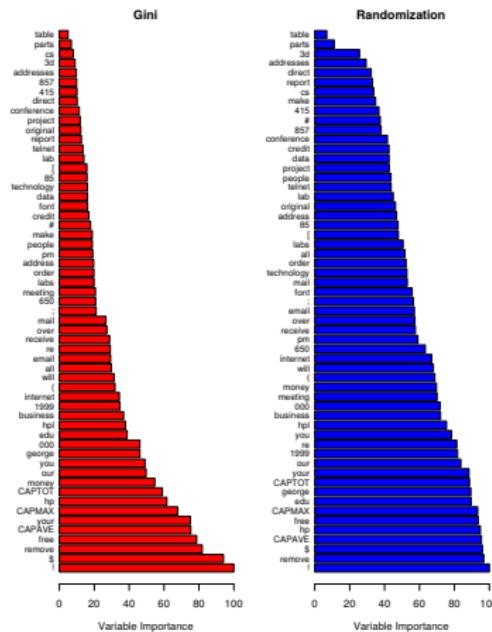
- Implemented by splitting the observations by considering only a *random* subset  $S \subset D$  of predictors.
- Helps to further reduce variance when predicted values are highly correlated across samples (due to correlated predictors).
- In bagging,  $S = D$ .
- Recommended values of  $S$ :  $S = \lfloor \frac{D}{3} \rfloor$  (regression);  $S = \sqrt{D}$  (classification)

## Algorithm

- ① For  $b = 1$  to  $B$ :
  - a Draw a bootstrap sample  $\mathcal{D}^*$  of size  $N$  from training dataset
  - b Grow a random-forest tree  $T_b$  to fit the sample by RBP as follows until minimum node size is reached:
    - i Randomly select subset  $S_m$  of the  $D$  predictors at node  $m$
    - ii Pick the best predictor/split-point  $(j, t)$  from  $S_m$
    - iii Split node  $m$  into two daughter nodes
- ② Output the ensemble of trees  $\{T_b\}_1^B$

# Feature importance in random forests

- Since the feature subset in each step of partitioning is random in all the  $B$  trees, the standard variable importance approach is not effective in random forests
- Instead, OOB samples are used to determine the variable importance
- This is done by perturbing each variable  $j$  in the OOB samples and then averaging the decrease in accuracy across the trees
- Illustrated on the spam dataset:



**Figure:** Variable importance computed on a scale of 0-100 from (Left) gradient boosted model (via Gini index) and (Right) random forest model (via OOB). Rankings are comparable, but relative values are not.

# Boosting

Generally, boosting refers to the approach of sequentially weighting the predictions of weak learners to obtain a final prediction

It can be considered as a method of fitting an [additive model](#):

$$f(\mathbf{x}; \boldsymbol{\theta}) = \sum_{m=1}^M \beta_m b_m(\mathbf{x}; \boldsymbol{\theta}_m) \quad (18)$$

where:

- $\beta_m$  are the expansion coefficients (weights)
- $b_m(\mathbf{x}; \boldsymbol{\theta}_m)$  are elementary (simple) basis functions of  $\mathbf{x}$  characterized by a set of parameters  $\boldsymbol{\theta}_m$ , e.g.
  - Shallow decision trees:  $\boldsymbol{\theta}_m$  parameterizes the split variable/split-point  $(j, t)$  at the internal nodes and the predictions at the terminal nodes
  - Single-hidden-layer neural networks:  $b(\mathbf{x}; \boldsymbol{\theta}_m) = \sigma(\mathbf{w}^T \mathbf{x})$ , where  $\sigma(t) = 1/(1 + e^{-t})$

# Boosting (cont.)

- To fit such additive models, we minimize the sum of the loss functions across the basis expansions.

$$\min_{\{\beta_m, \theta_m\}_1^M} \sum_{i=1}^N \ell \left( y_i, \sum_{m=1}^M \beta_m b(\mathbf{x}_i; \theta_m) \right) \quad (19)$$

- Choices for  $\ell$  include: squared error, likelihood-based, exponential (Adaboost)
- Fitting can be efficiently done via the *forward stagewise additive modeling* algorithm<sup>2</sup>

---

<sup>2</sup>ESL 10.1–10.5

# Forward stagewise additive modeling

① Initialize  $f_0(\mathbf{x}) = 0$

② For  $m = 1$  to  $M$ :

    a Find:

$$(\beta_m, \boldsymbol{\theta}_m) = \arg \min_{\beta, \boldsymbol{\theta}} \sum_{i=1}^M \ell(y_i, f_{m-1}(\mathbf{x}_i) + \beta b(\mathbf{x}_i; \boldsymbol{\theta}))$$

    b Set  $f_m(\mathbf{x}) = f_{m-1}(\mathbf{x}) + \beta_m b(\mathbf{x}; \boldsymbol{\theta}_m)$

If  $\ell$  is specified as squared error (typical for regression), then:

$$\ell(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2 \quad (20)$$

$$\ell(y_i, f_{m-1}(\mathbf{x}_i) + \beta b(\mathbf{x}_i; \boldsymbol{\theta})) = (y_i - f_{m-1}(\mathbf{x}_i) - \beta b(\mathbf{x}_i; \boldsymbol{\theta}))^2 \quad (21)$$

$$= (r_{im} - \beta b(\mathbf{x}_i; \boldsymbol{\theta}))^2 \quad (22)$$

where  $r_{im}$  is the **residual** of model  $m$  on the  $i$ th observation.

# Boosting trees

Boosting can be used as a “slow learning” approach for decision trees

- First we generate  $B$  bootstrap samples
- Starting with an initial tree, we fit a subsequent tree to the residuals from the first model
- We then update the tree and its residuals using a fraction  $\lambda$  of the predictions of the new tree
- The estimation and residual update step is repeated until the  $B$ th sample is used
- The boosted model is then given by the weighted sum of the fitted tree estimates

# Boosted regression tree algorithm

- ① Set  $\hat{f}(\mathbf{x}) = 0$  and  $r_i = y_i$  for all  $i$  in training dataset
- ② For  $b = 1, 2, \dots, B$ , repeat:
  - a Fit a tree  $\hat{f}^b$  with  $d$  splits ( $d + 1$  terminal nodes) to training data  $(\mathbf{x}, r)$
  - b Update  $\hat{f}$  by adding in a shrunken version of new tree:

$$\hat{f}(\mathbf{x}) \leftarrow \hat{f}(\mathbf{x}) + \lambda \hat{f}^b(\mathbf{x}) \quad (23)$$

- ③ c Update residuals:

$$r_i \leftarrow r_i - \lambda \hat{f}^b(\mathbf{x}_i) \quad (24)$$

- ③ Output boosted model:

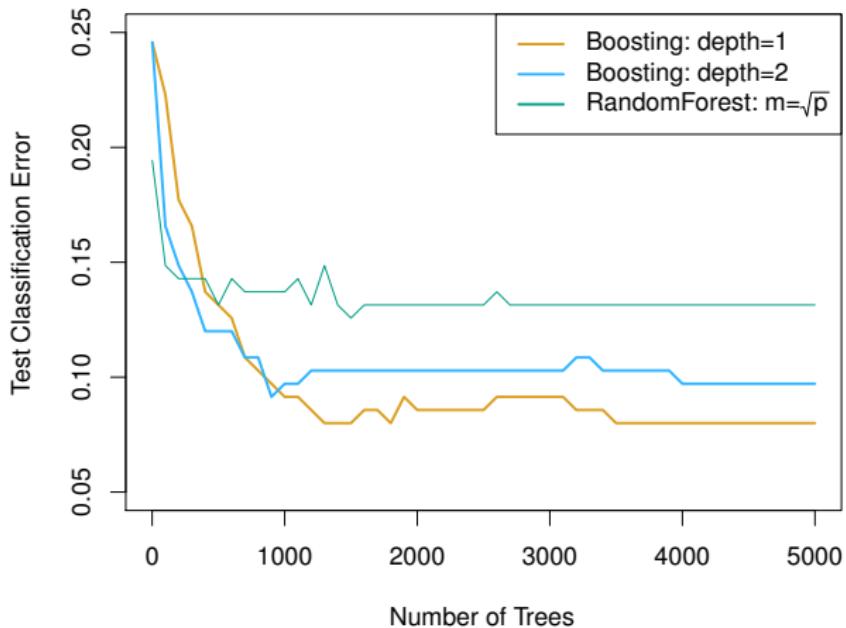
$$\hat{f}(\mathbf{x}) = \sum_{b=1}^B \lambda \hat{f}^b(\mathbf{x}) \quad (25)$$

# Notes on boosting trees

- There are three tuning parameters to consider:
  - Number of trees  $B$ . If  $B$  is too large, model is prone to overfitting. We choose  $B$  by cross-validation
  - Shrinkage/learning parameter  $\lambda > 0$  (typically 0.01 or 0.001)
  - Number of splits  $d$  in each tree (interaction depth);  $d = 1$  is a typical choice (in this case, the trees are referred to as “stumps”)
- The sequential fitting of  $\hat{f}(\mathbf{x})$  is analogous to *gradient descent* (since we are minimizing a loss function in each split)
  - Thus,  $r_i$  can be updated via the gradient of the loss function
  - Hence, the procedure is referred to as **gradient boosting**
  - The popular modeling package XGBoost (Extreme Gradient Boosting) is based on this idea but introduces regularization and second-order gradients

# Boosting performance

Boosting typically outperforms random forests and depth-1 trees perform better than other choices for  $d$ .



# Further issues and topics

- The methods described in this lecture are generally referred to as the CART implementation (classification and regression tree).
- Other approaches exist, e.g. PRIM (Patient Rule Induction Method) and ID3 (and its successors)

## Potential issues with trees and mitigations

- **Multilevel categorical variables:** the greater the number of levels, the more prone to the tree is to overfitting
- **Missing data:** dropping incomplete observations can reduce accuracy; imputation methods can be used
- **Tree instability:** errors or changes to data can drastically affect results  
*Bagging* can be used to address this (next lecture)
- **Lack of smoothness:** can be addressed using MARS (Multivariate Adaptive Regression Splines)

# Reading

- **PMLI** 18
- **ESL** Section 9.2
- **ISLR** 8.3
- Bagging: **ESL** 8.7
- Random Forests: **ESL** 15
- Boosting: **ESL** 10 (Very dense chapter. You may want to focus on 10.1–10.3, 10.7, 10.10, 10.13, 10.14.)

Further reading/study:

- **PRIM**: **ESL** 9.3 and Friedman's paper at  
<http://statweb.stanford.edu/~jhf/ftp/prim.pdf>.  
PRIM/CART can be used in scenario discovery. For a brief overview, see Jan Kwakkel's blog [post](#).
- **MARS**: **ESL** 9.4. This article provides a readily digestible overview:  
<http://uc-r.github.io/mars>.